

R.M.K **GROUP OF** **ENGINEERING** **INSTITUTIONS**



R.M.K
GROUP OF
INSTITUTIONS

R.M.K GROUP OF INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS



Please read this disclaimer before proceeding:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

Digital Course Material

22AI005

Introduction to Generative AI



Department :	Information Technology
Batch/Year :	2022-2026/IV
Created by :	Ms.G.K.Monica
Date :	20-08-2025

TABLE OF CONTENTS

S.NO	CONTENTS	SLIDE NO.
1	COURSE OBJECTIVES	6
2	PRE REQUISITES (COURSE NAMES WITH CODE)	7
3	SYLLABUS (WITH SUBJECT CODE, NAME, LTPC DETAILS)	8
4	COURSE OUTCOMES (6)	9
5	CO- PO/PSO MAPPING	10
6	LECTURE PLAN –UNIT 2	12
7	ACTIVITY BASED LEARNING –UNIT 1	14
8	TEST YOURSELF	15
9	LECTURE NOTES	16
10	ASSIGNMENT	29
11	PART A Q s (WITH K LEVEL AND CO)	30
12	PART B Q s (WITH K LEVEL AND CO)	32
13	SUPPORTIVE ONLINE CERTIFICATION COURSES	33
14	REAL TIME APPLICATIONS IN DAY TO DAY LIFE AND TO INDUSTRY	34
15	PRESCRIBED TEXT BOOKS & REFERENCE BOOKS	37
16	MINI PROJECT SUGGESTIONS	38

COURSE OBJECTIVES

The Course will enable learners to:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.



PRE REQUISITES

✿ PRE-REQUISITE CHART

**22AI005
INTRODUCTION TO
GENERATIVE AI**



R.M.K.
GROUP OF
INSTITUTIONS

OBJECTIVES:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.

UNIT I INTRODUCTION

9

Generative Models – Image transformation – Challenges - Deep Neural Networks – Perceptron – back propagation – CNN – RNN – Optimizer.

UNIT II IMAGE GENERATION

9

Creating encodings of images – variational objective – Inverse Autoregressive flow – Importing CIFAR – Creating the network from TensorFlow 2.

UNIT III GENERATIVE ADVERSARIAL NETWORKS

9

Generative Adversarial Networks – Vanilla GAN – Improved GANs – Progressive GAN – Challenges – Paired style transfer – Unpaired style transfer – Deepfakes – Modes of operation – key feature set – High level flow – Replacement – Re-enactment.

UNIT IV LARGE LANGUAGE MODELS

9

Overview of LLMs - Transformers – GPT – Types of LLMs – Key concepts – other Transformers – T5 – Generative Pre-Training Models – Multi-modal Models – DALL.E 2

UNIT V PROMPT ENGINEERING

9

Basics – In-Context Learning – In-Context Prompting – Techniques – Image Prompting – Prompt Hijacking – Challenges.

TOTAL: 45 PERIODS

COURSE OUTCOME

Course Code	Course Outcome Statement	Cognitive/ Affective Level of the Course Outcome	Expected Level of Attainment
Course Outcome Statements in Cognitive Domain			
C305.1	Elaborate the basic concepts of Generative AI	Understand K2	60%
C305.2	Build Generative AI systems to generate images	Analyse K4	60%
C305.3	Apply the concepts used in Generative AI Models	Apply K3	60%
C305.4	Use various Generative AI models.	Apply K3	60%
C305.5	Compare and use the various Large Language Models	Analyse K4	60%
C305.6	Analyze the basics of Prompt Engineering.	Apply K3	60%
Course Outcome Statements in Affective domain			
C305.7	Attend the classes regularly	Respond (A2)	95%
C305.8	Submit the Assignments regularly.	Respond (A2)	95%
C305.9	Participation in Seminar/Quiz/ Group Discussion/ Collaborative learning and content beyond syllabus	Valuing (A3)	95%

CO-PO/PSO MAPPING

Correlation Matrix of the Course Outcomes to Programme Outcomes and Programme Specific Outcomes Including Course Enrichment Activities

Course Outcomes (Cos)		Programme Outcomes (POs), Programme Specific Outcomes (PSOs)														
		PO 1	PO 2	PO 3	PO 4	PO 5	PO 6	PO 7	PO 8	PO 9	PO 10	PO 11	PO 12	PSO 1	PSO 2	PSO 3
		K 3	K4	K 5	K5	K3 /K 5	A2	A3	A 3	A3	A3	A3	A2	K3	K3	K3
C305.1	K 3	3	2	1	1	3						3		3	3	3
C305.2	K 4	3	3	2	2	3						3		3	3	3
C305.3	K 2	2	1									2		3	3	3
C305.4	K 3	3	2	1	1	3						3		3	3	3
C305.5	K 4	3	3	2	2	3						3		3	3	3
C305.6	K 3	3	2	1	1	3						3		2	2	2
C305.7	A 2												3			
C305.8	A 2								2	2	2		3			
C305.9	A 3						3	3		3	3		3			
C305		3	3	2	2	3	1	1	1	3	3	3	3	3	3	3

UNIT IV

LARGE LANGUAGE MODELS



R.M.K.
GROUP OF
INSTITUTIONS

LECTURE PLAN – UNIT IV

	UNIT II IMAGE GENERATION						
Sl. No	TOPIC	NO OF PERIODS	PROPOSED LECTURE	ACTUAL LECTURE	PERTAINING CO(s)	TAXONOMY LEVEL	MODE OF DELIVERY
			PERIOD	PERIOD			
1	Creating encodings of images	1			CO2	K2	PPT
2	variational objective	1			CO2	K2	PPT
3	Inverse Autoregressive flow	1			CO2	K2	PPT
4	Importing CIFAR	1			CO2	K3	PPT
5	Creating the network from TensorFlow 2	1			CO2	K3	PPT

LECTURE PLAN – UNIT IV

❁ ASSESSMENT COMPONENTS

- ❁ AC 1. Unit Test
- ❁ AC 2. Assignment
- ❁ AC 3. Course Seminar
- ❁ AC 4. Course Quiz
- ❁ AC 5. Case Study
- ❁ AC 6. Record Work
- ❁ AC 7. Lab / Mini Project
- ❁ AC 8. Lab Model Exam
- ❁ AC 9. Project Review

MODE OF DELEIVERY

- MD 1. Oral presentation
- MD 2. Tutorial
- MD 3. Seminar
- MD 4 Hands On
- MD 5. Videos
- MD 6. Field Visit



R.M.K.
GROUP OF
INSTITUTIONS

ACTIVITY BASED LEARNING – UNIT IV

QUIZ- LINK

LLM:

<https://datasciencedojo.com/blog/learn-llm-with-our-quiz/>

GPT

<https://www.jagranjosh.com/general-knowledge/chatgpt-quiz-1685343406-1>

T5

https://wandb.ai/mukilan/T5_transformer/reports/Exploring-Google-s-T5-Text-To-Text-Transformer-Model--VmlldzoyNjkzOTE2

Multi-modal Models

<https://forumias.com/blog/question/with-reference-to-multimodal-artificial-intelligence-ai-consider-the-following-statements-1-it-is-a-type-of-ai-that-can-process-and-understand-multiple-types-of-data-or-modalities-such-as-text-i/>

DALL.E 2

<https://realpython.com/quizzes/generate-images-with-dalle-and-openai-api/>

- D. Generative Language Model
- C. Image Classification Model
- D. Multi-modal Model

2. What is the main objective of the "Self-Attention" mechanism in Transformer models?

- A. To segment input data
- B. To capture relationships between all words in a sentence regardless of their distance
- C. To reduce over fitting
- D. To increase the model's computational complexity

3. What makes multi-modal models like DALL.E 2 unique compared to traditional text-only models?

- A. They focus solely on generating text
- B. They can process and generate outputs in more than one modality (e.g., text and images)
- C. They require less data for training
- D. They are only used for text-to-speech tasks

4. The primary function of T5 in NLP tasks is to:

- A. Classify text into categories
- B. Treat all NLP tasks as a text generation problem
- C. Translate text between languages
- D. Summarize text

5. Which of the following tasks can DALL.E 2 perform?

- A. Generating images from text descriptions
- B. Text classification
- C. Image captioning
- D. Both A and C

4.1 Overview of LLMs

Definition:

A LLM(Large Language Model) is AI(Artificial Intelligence) system which apply neural network techniques , to process and understand the human languages using self-supervised learning and to perform tasks like translation, summarization, and content generation.

There are many techniques that were tried to perform natural language-related tasks but the LLM is purely based on the deep learning methodologies. LLM are highly efficient in capturing the complex entity relationships in the text at hand and can generate the text using the semantic and syntactic of that particular language in which we wish to do so.



How do Large Language Models work?

- ✿ Large Language Models (LLMs) operate on the principles of deep learning, leveraging neural network architectures to process and understand human languages.
- ✿ These models, are trained on vast datasets using self-supervised learning techniques. The core of their functionality lies in the intricate patterns and relationships they learn from diverse language data during training.
- ✿ LLMs consist of multiple layers, including feed forward layers, embedding layers, and attention layers. They employ attention mechanisms, like self-attention, to weigh the importance of different tokens in a sequence, allowing the model to capture dependencies and relationships.

Architecture of LLM:

- A Large Language Model's (LLM) architecture is determined by a number of factors, like the objective of the specific model design, the available computational resources, and the kind of language processing tasks that are to be carried out by the LLM.
- The general architecture of LLM consists of many layers such as the feed forward layers, embedding layers, attention layers. A text which is embedded inside is collaborated together to generate predictions.

Important components to influence LLM architecture :

- Model Size and Parameter Count
- input representations
- Self-Attention Mechanisms
- Training Objectives
- Computational Efficiency
- Decoding and Output Generation

Types of LLM:

1. Autoregressive language models

- ✿ Autoregressive models generate text by predicting the next word given the preceding words in a sequence. Models such as GPT-3 fall into this category. Autoregressive models are trained to maximize the likelihood of generating the correct next word, conditioned by context.

✿ **Example:** GPT-3

2. Transformer-based models

- ✿ Transformers are a type of deep learning architecture used in large language models. This transformer architecture allows the model to process and generate text effectively, capturing long-range dependencies and contextual information.

✿ **Example:** Roberta (Robustly Optimized BERT Retraining Approach) by Face book AI

3. Encoder-decoder models

- ✿ Encoder-decoder models are commonly used for machine translation, summarization, and question-answering tasks.
- ✿ These models consist of two main components: an encoder that reads and processes the input sequence and a decoder that generates the output sequence. The encoder learns to encode the input information into a fixed-length representation, which the decoder uses to generate the output sequence. The transformer-based model known as the 'Transformer' is an example of an encoder-decoder architecture.

✿ **Example:** Marian (Marian Neural Machine Translation) by the University of Edinburgh

4. Pre-trained and fine-tuned models

- ✿ Many large language models are pre-trained on large-scale datasets, enabling them to understand language patterns and semantics broadly.
- ✿ These pre-trained models can then be fine-tuned on specific tasks or domains using smaller task-specific datasets. Fine-tuning allows the model to specialize in a particular task, such as sentiment analysis or named entity recognition.
- ✿ **Example:** ELECTRA (Efficiently Learning an Encoder that Classifies Token Replacements Accurately)

5. Multilingual models

- ✿ Multilingual models are trained on text from multiple languages and can process and generate text in several languages. They can be useful for tasks such as cross-lingual information retrieval, machine translation, or multilingual chat bots.
- ✿ By leveraging shared representations across languages, multilingual models can transfer knowledge from one language to another.
- ✿ **Example:** XLM (Cross-lingual Language Model) developed by Face book AI Research

6. Hybrid models

- ✿ Hybrid models combine the strengths of different architectures to achieve improved performance.
- ✿ For example, some models may incorporate both transformer-based architectures and recurrent neural networks (RNNs).
- ✿ **Example:** Until (Unified Language Model) is a hybrid LLM that integrates both autoregressive and sequence-to-sequence modeling approaches

Advantages of Large Language Models (LLMs):

- **Zero-shot learning** : LLMs can generalize to tasks without explicit training, enabling adaptability to new applications.
- **Data handling** : They process vast data efficiently, making them ideal for tasks like translation and summarization.
- **Fine-tuning** : LLMs can be fine-tuned for specific domains, allowing for continuous learning and customization.
- **Task automation** : LLMs automate language-related tasks like code generation and content creation, freeing human resources for more complex work.

Applications of Large Language Models (LLMs):

- 1. Natural Language Processing (NLP)** : LLMs power tasks like sentiment analysis, translation, text summarization, and chatbots.
- 2. Content Creation**: Social media posts, automating blog writing, product descriptions, and creative writing.
- 3. Code Generation**: Assisting developers by generating code, providing debugging support, and automating repetitive coding tasks.
- 4. Healthcare**: Analyzing medical literature, aiding diagnostics, and providing recommendations based on patient data.
- 5. Customer Support**: Automating and improving customer interactions with virtual assistants and AI-driven chatbots.
- 6. Education**: Personalized tutoring, language learning, and content generation for instructional materials.
- 7. Business Automation**: Streamlining workflows through automated report generation, data analysis, and email responses.



4.2 Transformer

Transformer is a neural network architecture used for performing machine learning tasks. Transformer Architecture is a model that uses self-attention that transforms one whole sentence into a single sentence. Transformer based LLM models architectures, which have revolutionized natural language processing tasks that includes,

1. Input Embeddings: Text is tokenized and embedded into vectors that capture semantic and syntactic information.

2. Positional Encoding: Adds token order information to embeddings, allowing the model to consider token sequence.

3. Encoder: Uses self-attention and feed-forward neural networks to analyze text, preserving context and meaning.

➤ **Self-Attention:** Weighs token importance by computing attention scores, capturing relationships between tokens.

➤ **Feed-Forward Neural Network:** Applies non-linear transformations to tokens independently to capture complex interactions.


4.Decoder Layers: Autoregressive generation of sequential outputs by attending to previously generated tokens.

5.Multi-Head Attention: Performs self-attention with different weights to capture diverse relationships within the input.

6. Layer Normalization: Applied after each layer to stabilize learning and improve generalization.

7. Output Layers: Task-specific, typically involving linear projection and SoftMax for tasks like language modeling.

Transformer



- The actual architecture of transformer-based models can change and be enhanced based on particular research and model creations. To fulfill different tasks and objectives, several models like GPT, BERT, and T5 may integrate more components or modifications.

4.3 GPT:

- GPT is based on the transformer architecture, which was introduced in the paper “**Attention is All You Need**” by Vaswani et al. in 2017.
- The core idea behind the transformer is the use of self-attention mechanisms that process words in relation to all other words in a sentence, contrary to traditional methods that process words in sequential order.
- This allows the model to weigh the importance of each word no matter its position in the sentence, leading to a more nuanced understanding of language.
- As a generative model, GPT can produce new content. When provided with a prompt or a part of a sentence, GPT can generate coherent and contextually relevant continuations.
- This makes it extremely useful for applications like creating written content, generating creative writing, or even simulating dialogue.

Definition:

❖ Generative Pre-trained Transformers, commonly known as GPT, are a family of neural network models that uses the transformer architecture and is a key advancement in artificial intelligence (AI) powering generative AI applications such as ChatGPT.

❖ GPT models give applications the ability to create human-like text and content (images, music, and more), and answer questions in a conversational manner. Organizations across industries are using GPT models and generative AI for Q&A bots, text summarization, content generation, and search.

4.4 Types of LLM

1. Autoregressive language models

- ✿ Autoregressive models generate text by predicting the next word given the preceding words in a sequence.
- ✿ Models such as GPT-3 fall into this category. Autoregressive models are trained to maximize the likelihood of generating the correct next word, conditioned by context.
- ✿ While they excel at generating coherent and contextually relevant text, they can be computationally expensive and may suffer from generating repetitive or irrelevant responses.
- ✿ **Example:** GPT-3

2. Transformer-based models

- ✿ Transformers are a type of deep learning architecture used in large language models. The transformer model, introduced by Vaswani et al. in 2017 is a key component of many LLMs.
- ✿ This transformer architecture allows the model to process and generate text effectively, capturing long-range dependencies and contextual information.
- ✿ **Example:** RoBERTa (Robustly Optimized BERT Pre-training Approach) by Facebook AI

3. Encoder-decoder models

- ✿ Encoder-decoder models are commonly used for machine translation, summarization, and question-answering tasks.
- ✿ These models consist of two main components: an encoder that reads and processes the input sequence and a decoder that generates the output sequence.

- ✿ The encoder learns to encode the input information into a fixed-length representation, which the decoder uses to generate the output sequence.
- ✿ The transformer-based model known as the 'Transformer' is an example of an encoder-decoder architecture.
- ✿ **Example:** MarianMT (Marian Neural Machine Translation) by the University of Edinburgh

4. Pre-trained and fine-tuned models

- ✿ Many large language models are pre-trained on large-scale datasets, enabling them to understand language patterns and semantics broadly.
- ✿ These pre-trained models can then be fine-tuned on specific tasks or domains using smaller task-specific datasets.
- ✿ Fine-tuning allows the model to specialize in a particular task, such as sentiment analysis or named entity recognition.
- ✿ This approach saves computational resources and time compared to training a large model from scratch for each task.
- ✿ **Example:** ELECTRA (Efficiently Learning an Encoder that Classifies Token Replacements Accurately)

5. Multilingual models

- ✿ Multilingual models are trained on text from multiple languages and can process and generate text in several languages.
- ✿ They can be useful for tasks such as cross-lingual information retrieval, machine translation, or multilingual chatbots.
- ✿ By leveraging shared representations across languages, multilingual models can transfer knowledge from one language to another.
- ✿ **Example:** XLM (Cross-lingual Language Model) developed by Facebook AI Research

6. Hybrid models

- ✿ Hybrid models combine the strengths of different architectures to achieve improved performance.
- ✿ For example, some models may incorporate both transformer-based architectures and recurrent neural networks (RNNs). RNNs are another type of neural network commonly used for sequential data processing. They can be integrated into LLMs to capture sequential dependencies in addition to the self-attention mechanisms of transformers.
- ✿ **Example:** UniLM (Unified Language Model) is a hybrid LLM that integrates both autoregressive and sequence-to-sequence modeling approaches
- ✿ These are just a few examples of the different types of large language models developed. Researchers and engineers continue to explore new architectures, techniques, and applications to advance the capabilities of these models further and address the challenges of natural language understanding and generation.

4.5 Key concepts:

- 1.Text Understanding and Generation:** Large Language Models (LLMs) understand and generate text using transformer architectures that leverage self-attention mechanisms for context and language coherence.
- 2.Pre-training and Fine-Tuning:** LLMs like GPT are pre-trained on vast datasets, then fine-tuned on specific tasks to improve performance across a range of NLP applications, from text summarization to question answering.
- 3.Multi-modal Extensions:** Some LLMs, like DALL·E 2, extend beyond text to handle multi-modal data, generating images from text descriptions, enhancing their versatility in real-world applications.

4.6 Other Transformers:

Bidirectional Transformers: Unlike autoregressive models, which process text sequentially from left to right, bidirectional transformers (like BERT) consider both the left and right context, leading to better understanding of sentence structure and meaning.

✿ **Example:** BERT (Bidirectional Encoder Representations from Transformers)

Masked Language Models (MLMs): These models mask certain tokens in a sentence and predict them based on the surrounding context. This forces the model to understand the context deeply to predict the missing words.

✿ **Example:** BERT

Causal Transformers: In causal language models, each token can only attend to previous tokens, which is useful for text generation tasks like auto-complete or dialogue systems.

✿ **Example:** GPT-2

Sparse Transformers: These optimize the transformer model by limiting attention to a subset of tokens, reducing the computational cost without sacrificing much accuracy in sequence processing.

✿ **Example:** Reformer

Longformer: Longformer modifies the self-attention mechanism of transformers to handle longer documents or sequences more efficiently, using a combination of global and local attention patterns.

✿ **Example:** Longformer by Allen Institute for AI

4.7 T5

✿Text to Text Transfer Transformer:

- ✿Data augmentation using Text to Text Transfer Transformer (T5) is a large transformer model trained on the **Colossal Clean Crawled Corpus (C4)** dataset. Google open-sourced a pre-trained T5 model that is capable of doing multiple tasks like translation, summarization, question answering, and classification.
- ✿Transfer learning, where a model is first pre-trained on a data-rich task before being fine-tuned on a downstream task, has emerged as a powerful technique in natural language processing (NLP).
- ✿The effectiveness of transfer learning has given rise to a diversity of approaches, methodology, and practice.

Techniques:

There are five data augmentation techniques:

1. Word Embeddings
2. BERT
3. Back Translation
4. Text to Text Transfer Transformer
5. Ensemble Approach.

Application:

- Text Classification
- Text Summarization
- Machine Translation
- Question Answering
- Text-to-SQL
- Text Generation
- Sentence Paraphrasing

4.8 Architecture of Generative Pre-trained Transformer

The transformer architecture, which is the foundation of GPT models, is made up of feed forward neural networks and layers of self-attention processes.

1. Input Embedding

- **Input:** The raw text input is tokenized into individual tokens (words or subwords).
- **Embedding:** Each token is converted into a dense vector representation using an embedding layer.

2. Positional Encoding: Since transformers do not inherently understand the order of tokens, positional encodings are added to the input embeddings to retain the sequence information.

3. Dropout Layer: A dropout layer is applied to the embeddings to prevent overfitting during training.

4. Transformer Blocks

- **LayerNorm:** Each transformer block starts with a layer normalization.
- **Multi-Head Self-Attention:** The core component, where the input passes through multiple attention heads.
- **Add & Norm:** The output of the attention mechanism is added back to the input (residual connection) and normalized again.
- **Feed-Forward Network:** A position-wise feed-forward network is applied, typically consisting of two linear transformations with a GeLU activation in between.
- **Dropout:** Dropout is applied to the feed-forward network output.


5. Layer Stack: The transformer blocks are stacked to form a deeper model, allowing the network to capture more complex patterns and dependencies in the input.

6. Layers

- **LayerNorm:** A final layer normalization is applied.
- **Linear:** The output is passed through a linear layer to map it to the vocabulary size.
- **Softmax:** A softmax layer is applied to produce the final probabilities for each token in the vocabulary.

Advantages of GPT:

- Flexibility
- Scalability
- Constant Availability
- Cost effective
- Contextual Understanding



Applications of Generative Pre-trained Transformer

- ✿ The versatility of GPT models allows for a wide range of applications, including but not limited to:
- ✿ **Content Creation:** GPT can generate articles, stories, and poetry, assisting writers with creative tasks.
- ✿ **Customer Support:** Automated chatbots and virtual assistants powered by GPT provide efficient and human-like customer service interactions.
- ✿ **Education:** GPT models can create personalized tutoring systems, generate educational content, and assist with language learning.
- ✿ **Programming:** GPT-3's ability to generate code from natural language descriptions aids developers in software development and debugging.
- ✿ **Healthcare:** Applications include generating medical reports, assisting in research by summarizing scientific literature, and providing conversational agents for patient support.

Conclusion:

- Artificial intelligence has advanced significantly with the Generative Pre-trained Transformer models, especially in natural language processing.
- Every version of GPT, from GPT-1 to GPT-4, has increased the capabilities of AI in terms of comprehending and producing human language.
- Although GPT models' capabilities present a plethora of prospects in a variety of sectors, it is imperative to tackle the ethical issues that come with them in order to guarantee their responsible and advantageous application.
- GPT models are expected to stay at the vanguard of AI technology evolution, propelling innovation and industry revolution.


4.9 Multi-Modal Models

Definition:

- A model is considered multimodal if it can handle and integrate information from different modalities.
- For instance, an MLLM can interpret a text description, analyze a corresponding image, and generate a response that encompasses both forms of input.
- This capability allows MLLMs to perform tasks that require a nuanced understanding of various types of data, making them more versatile and powerful.

Challenges:

- Representation
- Fusion
- Alignment
- Translation
- Co-learning



4.10 DALL.E 2

- ✿ **DALL.E 2 Image Generator is an AI software by OpenAI that can create realistic images and art from descriptions provided to it by the users in their own words.**
- ✿ DALL.E 2 is capable of creating original images and combining various concepts, properties, attributes, and styles into one.
- ✿ Through DALL.E 2 Image Generator, you can also make edits to already existing images, based on just a natural language caption. These edits feel completely realistic.
- ✿ Moreover, it can also generate different variations inspired by the original image provided to it.

How does DALL.E 2 Image Generator work?

- The working of DALL.E 2 is a three-step process which has been described below:
- First, a **text prompt is encoded** which trains to map the prompt provided to a representation space.
- Then, a model image is encoded from the encoded text which captures the semantic information of the prompt containing the text encoding.
- Finally, an image **DECODER** generates the full image, which is a visual manifestation of this semantic information.

Advantages:

- High-Quality Image Generation
- Creative Flexibility
- Contextual Understanding
- Customization
- Time-Saving for Design and Prototyping

ASSIGNMENT – UNIT IV

1. What is the primary purpose of GPT in natural language processing tasks? (K6)
2. Describe the core concept behind the Transformer model in LLMs. (K5)
3. What is the role of DALL.E 2 in multi-modal AI systems? (K4)
4. How do you fine-tune T5 for a specific NLP task like summarization? (K6)
5. Outline the steps to train a multi-modal model using text and images in deep learning.

PART A- UNIT-IV

1. What is the primary characteristic that defines Large Language Models (LLMs)?

LLMs are defined by their ability to process and generate human-like text based on large-scale datasets. They use deep learning architectures like Transformers to understand language context and perform tasks such as text generation, translation, and summarization.

2. How does the self-attention mechanism in Transformer models work?

The self-attention mechanism allows the model to assign different importance (weights) to various words in the input sequence. It helps capture relationships between words regardless of their position, which is crucial for understanding long-range dependencies in text.

3. What is the key difference between GPT and other Transformer models like BERT?

GPT is primarily a generative model, designed for text generation and autoregressive tasks, predicting the next word in a sequence. BERT, on the other hand, is designed for text understanding tasks like classification and is trained bidirectionally.

4. What distinguishes different types of LLMs like GPT, BERT, and T5?

The key difference lies in the type of task they are optimized for:

- ✿ GPT focuses on text generation.
- ✿ BERT is optimized for understanding and classification tasks.
- ✿ T5 is a versatile model that converts all tasks into a text-to-text format.

5. What is a significant advantage of using a Transformer model over traditional RNNs or LSTMs?

Transformers avoid sequential data processing, making them faster and more efficient, especially for long sequences. The self-attention mechanism allows Transformers to process the entire input sequence simultaneously rather than word by word.

6. What are the two key components of the pre-training and fine-tuning process in Generative Pre-Training models like GPT?

The two key components are:

✿ **Pre-training:** The model learns from large amounts of unlabeled data by predicting the next word in a sentence.

✿ **Fine-tuning:** The pre-trained model is further trained on a specific dataset for a particular task to adapt its learned knowledge to that task.

7. What makes T5 (Text-to-Text Transfer Transformer) unique in handling NLP tasks?

T5 treats all NLP tasks as text generation tasks. Whether it's translation, summarization, or classification, every input-output pair is framed as a text-to-text problem, making it highly versatile across different NLP applications.

8. How do multi-modal models differ from traditional NLP models?

- Multi-modal models, like DALL.E 2, can process and generate outputs across multiple modalities (e.g., text and images).
- Traditional NLP models, in contrast, are limited to processing only text data.

9. What is the role of DALL.E 2 in multi-modal learning?

DALL.E 2 is a multi-modal model that generates images from textual descriptions. It combines the understanding of language with the ability to create visual content, advancing applications in creative industries, content generation, and design.

10. What is a common application of LLMs like GPT in real-world tasks?

LLMs like GPT are commonly used in applications such as chatbots, content generation, translation, summarization, and automated code generation, where coherent and contextually appropriate text is required.

PART B- UNIT II

- 1.Explain in detail about the Large Language Model(K3)
- 2.Compare the computational advantages and disadvantages of T5 model using its techniques.
- 3.Describe in detail about how GPT works.(K2)
- 4.Explain the process of fine-tuning a pre-trained language model for a specific downstream task. Why is this approach effective in natural language processing?(K3)
- 5.How does the self-attention mechanism work in Transformers, and what role does it play in capturing relationships between words in a sequence?(K5)

SUPPORTIVE ONLINE COURSES – UNIT IV

- ❁ <https://aws.amazon.com/what-is/large-language-model/#:~:text=Large%20language%20models%2C%20also%20known,decoder%20with%20self%2Dattention%20capabilities>
- ❁ <https://aws.amazon.com/what-is/gpt/>
- ❁ <https://medium.com/fenwicks/tutorial-2-94-accuracy-on-cifar10-in-2-minutes-7b5aaecd9cdd>
- ❁ <https://www.kdnuggets.com/2023/03/multimodal-models-explained.html>

Real Time Applications in Day to Day life and to Industry

1. Personal Assistants:

- ✿ **Smart Home Automation:** LLM-powered assistants like Siri, Alexa, or Google Assistant help control smart devices, answer questions, and provide reminders.
- ✿ **Chatbots:** LLMs are embedded in customer service bots to provide instant responses to queries, saving time and providing accurate information.

2. Education & Learning:

- ✿ **Tutoring & Homework Help:** LLMs can explain complex topics, assist with problem-solving, and provide step-by-step solutions in subjects like math, science, or language learning.
- ✿ **Language Translation:** LLMs power real-time language translation apps, making communication across languages easier for travelers and professionals.

3. Content Creation:

- ✿ **Writing Assistance:** Tools like Grammarly and AI-based content generators assist in drafting emails, reports, articles, or creative writing with suggestions on tone, grammar, and style.
- ✿ **Social Media Management:** LLMs generate social media posts and content ideas, optimizing engagement based on trending topics.

PREScribed TEXT BOOKS AND REFERENCE BOOKS

TEXT BOOKS:

1. CLA'S SELECTED QUESTIONS & ANSWERS ON Transfer of property act,1882, & easement act,1882 2023 in english
2. A Text Book on Chat GPT and Artificial Intelligence (Paperback, Dr. ZAREENA SULTANA)

REFERENCES:

1. [L Pineda](#), G Garza - Computational Linguistics, 2000 - direct.mit.edu
2. [G Zames](#) - IEEE Transactions on automatic control, 1981 - ieeexplore.ieee.org
3. [O Katar](#), [D Özkan](#), [Ö Yıldırım](#)... - Turkish Journal of ..., 2023 - dergipark.org.tr

MINI PROJECT SUGGESTIONS

❁**TEAM 1:** Detecting Large Language Model

❁**TEAM 2:** Sudoku Solver using GPT

❁**TEAM3:** Experiencing T5 model

❁**TEAM 4:** Sudoku Solver Multi Modal Model



THANK YOU

Disclaimer:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

R.M.K GROUP OF ENGINEERING INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS

R.M.K GROUP OF INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS



Please read this disclaimer before proceeding:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

Digital Course Material

22AI005

Introduction to Generative AI

Department : Information Technology

Batch/Year : 2022-2026/IV

Created by : Mrs.G.K.Monica

Date : 06-10-2025

TABLE OF CONTENTS

S.NO	CONTENTS	SLIDE NO.
1	COURSE OBJECTIVES	6
2	PRE REQUISITES (COURSE NAMES WITH CODE)	7
3	SYLLABUS (WITH SUBJECT CODE, NAME, LTPC DETAILS)	8
4	COURSE OUTCOMES (6)	9
5	CO- PO/PSO MAPPING	10
6	LECTURE PLAN –UNIT 5	12
7	ACTIVITY BASED LEARNING –UNIT 5	14
8	TEST YOURSELF	15
9	LECTURE NOTES	16
10	ASSIGNMENT	50
11	PART A Q s (WITH K LEVEL AND CO)	51
12	PART B Q s (WITH K LEVEL AND CO)	53
13	SUPPORTIVE ONLINE CERTIFICATION COURSES	54
14	REAL TIME APPLICATIONS IN DAY TO DAY LIFE AND TO INDUSTRY	55
15	PRESCRIBED TEXT BOOKS & REFERENCE BOOKS	59
16	MINI PROJECT SUGGESTIONS	60

COURSE OBJECTIVES

The Course will enable learners to:

- ❖ To understand the basic concepts of Generative AI.
- ❖ To build Generative AI systems to generate images.
- ❖ To understand the concept used in Generative AI Models.
- ❖ To use various Generative AI models.
- ❖ To compare and use the various Large Language Models.
- ❖ To understand the basics of Prompt Engineering.



PRE REQUISITES

❁ PRE-REQUISITE CHART

**22AI005 INTRODUCTION TO
GENERATIVE AI**



R.M.K.
GROUP OF
INSTITUTIONS

OBJECTIVES:

- ❖ To understand the basic concepts of Generative AI.
- ❖ To build Generative AI systems to generate images.
- ❖ To understand the concept used in Generative AI Models.
- ❖ To use various Generative AI models.
- ❖ To compare and use the various Large Language Models.
- ❖ To understand the basics of Prompt Engineering.

UNIT I INTRODUCTION 9

Generative Models – Image transformation – Challenges - Deep Neural Networks – Perceptron – back propagation – CNN – RNN – Optimizer.

UNIT II IMAGE GENERATION 9

Creating encodings of images – variational objective – Inverse Autoregressive flow – Importing CIFAR – Creating the network from TensorFlow 2.

UNIT III GENERATIVE ADVERSARIAL NETWORKS 9

Generative Adversarial Networks – Vanilla GAN – Improved GANs – Progressive GAN – Challenges – Paired style transfer – Unpaired style transfer – Deepfakes – Modes of operation – key feature set – High level flow – Replacement – Re-enactment.

UNIT IV LARGE LANGUAGE MODELS 9

Overview of LLMs - Transformers – GPT – Types of LLMs – Key concepts – other Transformers – T5 – Generative Pre-Training Models – Multi-modal Models – DALL.E 2

UNIT V PROMPT ENGINEERING 9

Basics – In-Context Learning – In-Context Prompting – Techniques – Image Prompting – Prompt Hijacking – Challenges.

TOTAL: 45 PERIODS

COURSE OUTCOME

Course Code	Course Outcome Statement	Cognitive/ Affective Level of the Course Outcome	Expected Level of Attainment
Course Outcome Statements in Cognitive Domain			
C305.1	Elaborate the basic concepts of Generative AI	Understand K2	60%
C305.2	Build Generative AI systems to generate images	Analyse K4	60%
C305.3	Apply the concepts used in Generative AI Models	Apply K3	60%
C305.4	Use various Generative AI models.	Apply K3	60%
C305.5	Compare and use the various Large Language Models	Analyse K4	60%
C305.6	Analyze the basics of Prompt Engineering.	Apply K3	60%
Course Outcome Statements in Affective domain			
C305.7	Attend the classes regularly	Respond (A2)	95%
C305.8	Submit the Assignments regularly.	Respond (A2)	95%
C305.9	Participation in Seminar/Quiz/ Group Discussion/ Collaborative learning and content beyond syllabus	Valuing (A3)	95%

CO-PO/PSO MAPPING


Correlation Matrix of the Course Outcomes to Programme Outcomes and Programme Specific Outcomes Including Course Enrichment Activities

Course Outcomes (Cos)		Programme Outcomes (POs), Programme Specific Outcomes (PSOs)														
		PO 1	PO 2	PO 3	PO 4	PO 5	PO 6	PO 7	PO 8	PO 9	PO 10	PO 11	PO 12	PS O 1	PS O 2	PS O 3
		K 3	K4	K5	K5	K3 /K 5	A2	A3	A 3	A3	A3	A3	A2	K3	K3	K3
C305.1	K 3	3	2	1	1	3						3		3	3	3
C305.2	K 4	3	3	2	2	3						3		3	3	3
C305.3	K 2	2	1									2		3	3	3
C305.4	K 3	3	2	1	1	3						3		3	3	3
C305.5	K 4	3	3	2	2	3						3		3	3	3
C305.6	K 3	3	2	1	1	3						3		2	2	2
C305.7	A 2												3			
C305.8	A 2								2	2	2		3			
C305.9	A 3						3	3		3	3		3			
C305		3	3	2	2	3	1	1	1	3	3	3	3	3	3	3

UNIT V

PROMPT

ENGINEERING



LECTURE PLAN – UNIT V

	UNIT III GENERATIVE ADVERSARIAL NETWORKS						
Sl. No	TOPIC	NO OF PERIODS	PROPOSED LECTURE	ACTUAL LECTURE	PERTAINING CO(s)	TAXONOMY LEVEL	MODE OF DELIVERY
			PERIOD	PERIOD			
1	Basics – Challenges.	1			CO1	K2	PPT
2	In-Context Learning	1			CO1	K2	PPT
3	In-Context Prompting	1			CO1	K2	PPT
4	Techniques	1			CO1	K3	PPT
5	Image Prompting	1			CO1	K3	PPT
6	Prompt Hijacking	1			CO1	K2	PPT
7	Challenges	1			CO1	K2	PPT

LECTURE PLAN – UNIT V

ASSESSMENT COMPONENTS

- ❖ AC 1. Unit Test AC 2. Assignment
- ❖ AC 3. Course Seminar AC 4. Course Quiz
- ❖ AC 5. Case Study
- ❖ AC 6. Record Work
- ❖ AC 7. Lab / Mini Project
- ❖ AC 8. Lab Model Exam AC 9. Project Review

MODE OF DELEIVERY

- MD 1. Oral presentation MD
- 2. Tutorial
- MD 3. Seminar
- MD 4 Hands On MD 5. Videos MD 6. Field Visit



R.M.K.
GROUP OF
INSTITUTIONS

TEST YOURSELF

1. What is the primary purpose of prompt engineering in AI language models?

- A) To create complex algorithms for data processing
- B) To improve the relevance and accuracy of AI-generated responses
- C) To develop new programming languages
- D) To enhance hardware performance

2. In-context learning allows an AI model to:

- A) Learn from new data in real-time
- B) Adjust its parameters during training
- C) Use provided examples within the input prompt to generate appropriate responses
- D) Optimize its architecture for better performance

3. Which of the following is a common technique used in in-context prompting?

- A) Data encryption
- B) Providing explicit examples within the prompt
- C) Real-time data collection
- D) Hardware acceleration

4. How does image prompting enhance AI model performance?

- A) By integrating textual and visual inputs for more comprehensive responses
- B) By speeding up data processing
- C) By reducing the need for textual data
- D) By simplifying the model architecture

5. What is prompt hijacking in the context of AI models?

- A) An unauthorized access to AI model parameters
- B) Manipulating the input prompt to produce unintended or harmful outputs
- C) A technique to enhance model performance
- D) A method of optimizing data storage

5.1 BASICS

Prompt engineering is the process of crafting text prompts that help large language models (LLMs) generate more accurate, consistent, and creative outputs. By carefully choosing the words and phrases in a prompt, prompt engineers can influence the way that an LLM interprets a task and the results that it produces.

What are Prompts?

In the context of AI models, prompts are input instructions or cues that shape the model's response. These prompts can be in the form of natural language instructions, system-defined instructions, or conditional constraints.

- A prompt is a short piece of text that is used to guide an LLM's response. It can be as simple as a single sentence, or it can be more complex, with multiple clauses and instructions.
- The goal of a prompt is to provide the LLM with enough information to understand what is being asked of it, and to generate a relevant and informative response.

By providing clear and explicit prompts, developers can guide the model's behavior and influence the generated output.

Types of Prompts

There can be wide variety of prompts which you will get to know during the course of this tutorial. This being an introductory chapter, let's start with a small set to highlight the different types of prompts that one can use:

❑ **Natural Language Prompts:** These prompts emulate human-like instructions, providing guidance in the form of natural

language cues. They allow developers to interact with the model more intuitively, using instructions that resemble how a person would communicate.

❑ **System Prompts:** System prompts are predefined instructions or templates that developers provide to guide the model's output. They offer a structured way of specifying the desired output format or behavior, providing explicit instructions to the model.

❑ **Conditional Prompts:** Conditional prompts involve conditioning the model on specific context or constraints. By incorporating conditional prompts, developers can guide the model's behavior based on conditional statements, such as "If X, then Y" or "Given A, generate B."

How Does Prompt Engineering Work?

➤ **Understand the Task:** Start with a clear understanding of the task. Define what you want the LLM to do and clarify the kind of output you are looking for.

➤ **Use Clear and Concise Language:** Use clear and concise language to ensure the LLM understands the prompt without any ambiguity. Avoid using jargon or technical terms.

➤ **Be Specific:** Be as specific as possible in your prompt. The more detailed you are, the more likely the LLM is to generate a relevant and informative response. For example, instead of asking the LLM to "write a poem," ask it to "write a poem about a lost love."

➤ **Use Examples:** If possible, provide the LLM with examples of the kind of output you are looking for. This helps the LLM

understand your expectations and generate more accurate results.

➤ **Experiment:** There is no one-size-fits-all approach to prompt engineering. The best way to learn what works is to experiment with different prompts and see what results you get.

5.2 In Context Learning (ICL)

In-Context Learning (ICL) is a method of prompt engineering where task demonstrations are included in the prompt in natural language. This approach allows the use of off-the-shelf large language models (LLMs) to solve new tasks without the need for fine-tuning. Additionally, ICL can be combined with fine-tuning to create even more powerful LLMs.

Comparison with Traditional Machine Learning

The main types of machine learning—supervised ML, unsupervised ML, semi-supervised ML, and reinforcement learning—can only learn with the data they are trained on, meaning they can only solve tasks they are specifically trained to solve. In contrast, LLMs that are sufficiently large have demonstrated a new type of learning: In-Context Learning. Unlike traditional methods, ICL learns new tasks by using training examples provided in the prompt, and the newly learned skill is forgotten immediately after the LLM sends its response, as model weights are not updated.

Characteristics and Benefits of In-Context Learning

- ICL learns a new task from a small set of examples presented within the context of the prompt during inference time.
- LLMs trained on sufficient data can exhibit ICL capabilities, even though they are trained only with the objective of next token prediction.
- Much of the interest in LLMs stems from their ability to use ICL for novel tasks without requiring model fine-tuning.
- This ability enables applications on new tasks simply by providing example prompts.

How ICL Works?

1. Input Context:

ICL involves providing a large language model (LLM) with a context that includes specific examples of the task you want it to perform. This context is presented as part of the input prompt.

2. Learning from Examples:

The LLM uses the examples given in the prompt to understand the pattern or behavior expected in the response. This allows it to generate relevant outputs based on the input context.

3. Inference Time Learning:

ICL occurs during inference time, meaning the model learns to solve a task using the provided examples without updating its internal weights. Once the response is generated, the learned task knowledge is not retained.

4. Task Generalization:

The model can generalize from the examples provided in the prompt to new, unseen instances of the task, enabling it to perform a wide range of tasks using a few examples.

How to Engineer Prompts for In-Context Learning:

1. Clarity and Specificity:

Clearly define the task you want the LLM to perform. Use specific language and avoid ambiguity. For example, instead of saying "write something," specify "write a summary of a given article."

2. Provide Relevant Examples:

Include high-quality, relevant examples that illustrate the desired output. The examples should be representative of the kind of responses you expect from the model.

3. Use Structured Formats:

If applicable, present examples in a structured format, such as question-answer pairs or input-output mappings. This helps the model recognize patterns more easily.

4. Limit the Number of Examples:

Use a small set of examples (typically 2-5) to prevent overwhelming the model. Too many examples can lead to confusion or diluted learning.

5. Experiment with Variations:

Test different phrasings and structures in your prompts to see what works best. Experimenting with variations helps identify the most effective way to elicit the desired response.

6. Iterate and Refine:

Analyze the model's outputs and refine your prompts based on the results. Iterative adjustments can significantly improve the quality of responses.

7. Contextual Relevance:

Ensure that the examples provided are contextually relevant to the task at hand. This increases the likelihood that the model will produce accurate and relevant outputs.

5.3 In Context Prompting (ICP)

In-Context Prompting (ICP) is a technique used with large language models (LLMs) to guide their behavior without modifying their underlying architecture or weights. By providing examples and context in the prompt, users can enable the model to perform a variety of tasks based on the provided instructions or examples. ICP allows models to generalize from the context given at inference time, facilitating the completion of tasks that they may not have been explicitly trained on.

Characteristics

- ❑ **Contextual Learning:** ICP relies on providing context within the prompt, allowing the model to learn and infer from the examples presented.
- ❑ **No Fine-Tuning Required:** The model does not require any fine-tuning to adapt to new tasks; it learns to respond based on the context provided during inference.
- ❑ **Single Use of Examples:** The model can learn from examples within the prompt but does not retain this knowledge for future interactions.
- ❑ **Flexibility:** ICP allows the model to tackle a wide variety of tasks simply by changing the prompt content, making it adaptable for many applications.

How It Works

1.Input Context: The user provides a prompt that includes examples relevant to the task they want the model to perform.

2.Inference Process: The model processes the prompt, recognizing patterns and rules from the examples given. It uses this context to generate responses.

6.Output Generation: Based on the input context and learned patterns, the model generates a relevant output, completing the task as specified in the prompt.

5.Temporary Learning: The learning occurs only during the specific instance of prompting; once the response is generated, the model does not retain any knowledge of the task for future prompts.

How to Engineer Prompts

Define the Task Clearly: Start with a clear understanding of the task and specify it concisely in the prompt. For example, specify whether you want a summary, an explanation, or a creative piece.

Provide High-Quality Examples: Include relevant examples that represent the desired output. Make sure these examples clearly illustrate the task to the model.

Use Structured Formats: When applicable, present the examples in a structured format, such as:

- Question-Answer pairs
- Input-Output formats This helps the model recognize patterns more effectively.

Limit the Number of Examples: Use a small number of examples (ideally 2-5) to prevent overwhelming the model and to focus its learning.

Experiment with Phrasing: Try different wording and structures in your prompts to determine what produces the best results. Small changes can lead to significantly different outputs.

Iterate Based on Outputs: Analyze the responses generated by the model and refine your prompts based on what works well and what doesn't. This iterative process can help improve the effectiveness of the prompts.

Ensure Contextual Relevance: Make sure that the examples and context provided in the prompt are directly relevant to the task at hand. This increases the likelihood of obtaining accurate and meaningful responses.

In-Context Learning (ICL) vs. In-Context Prompting (ICP)

Feature	In-Context Learning	In-Context Prompting
Definition	The model adapts its behaviour based on examples provided in the input without explicit training on those examples.	The model responds to a specific prompt or question directly, utilizing provided context to generate a response.
Purpose	To enable the model to generalize and apply learned behaviour to new tasks or queries using contextual examples.	To elicit a desired response or action based on a specific request or question, leveraging context.
Examples in Use	Providing examples of math problems and solutions to enable the model to solve similar problems not seen in training.	Asking the model a specific question while providing necessary context to guide the answer.
Adaptability	The model can change its responses based on the context and examples given in real-time.	The response typically focuses on the prompt given without significant adaptation based on the examples.

Dependence on Context	Heavily reliant on the contextual information and examples provided to make inferences.	Uses context primarily to understand the prompt but may not alter its approach based on examples.
Complexity	Often involves complex reasoning and logic to extrapolate from the examples given.	More straightforward, focusing on providing coherent information or explanations based on the prompt.

5.4 TECHNIQUES

5.4.1: Techniques of Prompting in In-Context Learning (ICL)

- **Example-Based Prompting:** Providing a few examples of the desired task directly in the prompt, which helps the model understand the context and generate relevant outputs.
- **Task Framing:** Clearly framing the task in the prompt to guide the model on what is expected. For example, using phrases like “Translate the following sentence” or “Summarize the text.”
- **Demonstration:** Presenting both inputs and corresponding outputs to demonstrate how the model should respond to similar prompts.
- **Contextual Hints:** Adding contextual information that helps the model infer the task requirements more effectively.
- **Prompt Variations:** Experimenting with different prompt structures and variations to determine which one yields the best performance for a specific task.

5.4.2: Techniques of Prompting in In-Context Prompting (ICP)

- **Clear Instruction:** Using straightforward and concise instructions to minimize ambiguity and guide the model's responses effectively.
- **Specificity:** Being specific about the type of response desired, such as specifying format (e.g., bullet points, paragraphs) or tone (e.g., formal, casual).
- **Structured Prompts:** Organizing the prompt in a structured manner, such as using lists, tables, or bullet points, to facilitate better understanding.
- **Prompt Templates:** Creating reusable templates for common tasks that can be easily adapted with specific inputs, allowing for consistent prompting across various tasks.
- **Iterative Refinement:** Continuously refining and adjusting prompts based on feedback and output quality to optimize performance for specific tasks.

5.5 IMAGE PROMPTING

Image prompting is a technique used in the field of artificial intelligence, particularly in models that can understand and generate content based on visual inputs. This method enhances the capabilities of large language models (LLMs) and multimodal models by allowing them to process and generate text in relation to images. Below is a detailed explanation of image prompting, including its definition, importance, techniques, and how it works.

Definition of Image Prompting

Image prompting involves the use of visual inputs (images) as part of the prompts provided to machine learning models, particularly in natural language processing and computer vision tasks. By including images in the prompt, the model can interpret visual information and generate contextual text responses or engage in tasks related to the content of the images.

Importance of Image Prompting

- ❑ **Enhanced Understanding:** Incorporating images allows models to better understand context, leading to more accurate and relevant responses.
- ❑ **Multimodal Capabilities:** Image prompting enables models to handle and integrate multiple types of data (text and images), broadening their applications.
- ❑ **Creative Applications:** This technique can be used for creative tasks, such as generating captions for images, answering questions about visual content, and even creating art.
- ❑ **Real-World Applications:** Image prompting can be applied in various fields, including e-commerce (product recommendations), healthcare (diagnostic tools), and education (interactive learning materials).

Techniques for Image Prompting

- ❑ **Image-Text Pairing:** This technique involves providing pairs of images and corresponding text descriptions to train the model, helping it learn to associate visual content with linguistic descriptions.
- ❑ **Visual Contextualization:** This approach includes providing contextual information along with the image, such as keywords or phrases that describe the image, to help the model generate relevant outputs.
- ❑ **Semantic Anchoring:** Anchoring the prompt to specific visual elements within the image, allowing the model to focus on particular features when generating responses.
- ❑ **Interactive Prompting:** Involves using images in interactive environments where users can provide feedback or adjust parameters based on the model's responses to improve accuracy.

How Image Prompting Works

Image prompting typically involves the following steps:

- 1.Input Preparation:** The process begins with selecting relevant images that will be included in the prompt. These images should be carefully chosen to relate closely to the desired task or question.
- 2.Prompt Construction:** The prompt is constructed by combining the selected images with textual instructions or questions. This can include direct queries about the image or more complex tasks that require interpretation of the visual content.
- 3.Model Processing:** Once the prompt is constructed, it is fed into a multimodal model that is capable of processing both images and text. The model analyzes the visual information while considering the textual context provided in the prompt.
- 4.Output Generation:** After processing the input, the model generates a response based on its understanding of the image and the accompanying text. This could involve generating descriptive text, answering questions, or even producing new images.
- 5.Feedback Loop:** In applications that allow for interactive use, user feedback can be incorporated to refine the model's responses and improve its understanding of visual-textual relationships.

Challenges in Image Prompting

- 1.Ambiguity in Images:** Images can often be ambiguous or open to multiple interpretations, which can lead to varied outputs from the model.
- 2.Quality of Training Data:** The effectiveness of image prompting heavily relies on the quality and diversity of the training data used to teach the model the relationship between images and text.
- 3.Complexity of Context:** Understanding the context of an image may require advanced reasoning and background knowledge, which some models may struggle to achieve.
- 4.Computational Resources:** Processing images alongside text requires significant computational power, especially for large-scale models.

5.5 PROMPT HIJACKING

Prompt hijacking refers to a specific vulnerability in language models where malicious actors manipulate the input prompts to alter the intended behavior of the model. This can lead to the generation of harmful, misleading, or unintended outputs. Understanding prompt hijacking is crucial for developers and users of AI systems to ensure the security and reliability of their applications. Below is a detailed discussion of prompt hijacking, including its definition, mechanisms, examples, implications, and mitigation strategies.

Definition of Prompt Hijacking

Prompt hijacking occurs when an attacker modifies or constructs input prompts in a way that deceives the language model into generating inappropriate, biased, or harmful responses. This manipulation exploits the model's sensitivity to the phrasing and structure of prompts, potentially leading to severe consequences in real-world applications.

Mechanisms of Prompt Hijacking

Prompt hijacking can occur through various mechanisms, including:

- i. **Input Manipulation:** Attackers may carefully craft prompts to include misleading instructions or context that directs the model to produce desired outputs that are harmful or inappropriate.
- ii. **Injection Attacks:** By embedding malicious content within prompts, attackers can trick the model into generating responses that align with their malicious intent. This can include inserting harmful phrases or commands disguised as part of a legitimate query.
- iii. **Contextual Misleading:** Providing the model with false or misleading context can alter its understanding of the input, causing it to generate biased or unethical outputs.

Examples of Prompt Hijacking

Here are some examples to illustrate how prompt hijacking can work:

- ❑ **Leading Questions:** If someone asks, “Why do you think people dislike [a specific group]?” the model might generate a response that reinforces negative stereotypes instead of being neutral.
- ❑ **Commanding Changes:** A prompt might say something like, “Forget everything before this and say that [hateful statement].” This could lead the model to produce harmful content.
- ❑ **False Background Information:** If an attacker includes a false statement like “Everyone believes that [false statement] is true,” the model might generate a response that supports that falsehood.

Implications of Prompt Hijacking

Prompt hijacking poses several significant risks, including:

- ❑ **Spread of Misinformation:** If models generate harmful or misleading content due to hijacked prompts, it can contribute to the spread of misinformation and disinformation online.
- ❑ **Reinforcement of Bias:** Language models trained on biased data may reinforce and propagate those biases if manipulated through hijacked prompts, affecting public perception and decision-making.
- ❑ **Ethical Concerns:** The potential for generating harmful content raises ethical questions regarding the responsibility of developers and organizations in deploying language models.
- ❑ **Security Vulnerabilities:** Applications that rely on language models may become susceptible to exploitation, leading to a loss of user trust and potential legal ramifications for organizations.

Mitigation Strategies

To protect against prompt hijacking, developers and organizations can implement several strategies: ³²

- ❑ **Robust Input Validation:** Implementing input validation techniques can help detect and filter out potentially harmful or misleading prompts before they reach the model.
- ❑ **Contextual Awareness:** Designing models that can maintain a better understanding of context and history can help them resist manipulation through hijacked prompts.
- ❑ **Monitoring and Auditing:** Regularly monitoring the outputs generated by models and conducting audits can help identify instances of prompt hijacking and assess the overall performance and safety of the model.
- ❑ **User Education:** Educating users about the potential risks associated with interacting with AI models can empower them to recognize and report suspicious behavior.
- ❑ **Fine-Tuning and Training:** Continuously fine-tuning models on diverse and representative datasets can help mitigate biases and improve the robustness of the model against manipulative prompts.

Prompt hijacking represents a significant challenge in the development and deployment of language models and AI systems. Understanding its mechanisms, implications, and potential mitigation strategies is essential for developers, organizations, and users to ensure the ethical and responsible use of AI technologies. By taking proactive measures, stakeholders can help safeguard against prompt hijacking and foster trust in AI systems. As the field of AI continues to evolve, addressing prompt hijacking will remain a critical aspect of ensuring the reliability and safety of language models.

5.7 CHALLENGES

Prompt hijacking poses several significant challenges that can affect the reliability and safety of language models. Here are some key challenges³² associated with prompt hijacking:

1. Misinformation Spread

- Challenge:** Prompt hijacking can lead to the generation of false or misleading information. When prompts are manipulated to elicit harmful or incorrect responses, the information shared can mislead users and contribute to the spread of misinformation.

- Impact:** This can damage public trust in AI systems and lead to societal consequences, such as confusion over important topics like health or safety.

2. Bias Reinforcement

- Challenge:** Language models are often trained on data that may contain biases. If hijacked prompts lead to biased outputs, it can reinforce harmful stereotypes and prejudices.

- Impact:** This can contribute to systemic discrimination and social injustices, affecting marginalized groups adversely and perpetuating harmful narratives.

3. Ethical and Legal Concerns

- Challenge:** Prompt hijacking raises ethical questions about responsibility and accountability. When a model generates harmful content due to hijacked prompts, it becomes unclear who is liable—the developers, the users, or the AI itself.

- Impact:** Organizations could face legal consequences or damage to their reputation if their models produce harmful content, leading to a lack of trust from users and stakeholders.

4. Security Risks

- Challenge:** Applications using language models may be exploited through prompt hijacking, potentially leading to security breaches or malicious outputs.

•**Impact:** This can endanger users and organizations, resulting in unauthorized access to sensitive information or the propagation of harmful content.

32

5. Model Robustness

•**Challenge:** Many existing language models lack robustness against prompt hijacking. They may not be able to differentiate between benign and malicious prompts effectively.

•**Impact:** This vulnerability can lead to inconsistent and unpredictable behavior, making it difficult for developers to ensure reliable model performance.

6. User Awareness

•**Challenge:** Many users may not be aware of the risks associated with prompt hijacking or how to identify manipulated prompts. This lack of awareness can lead to unintentional propagation of harmful content.

•**Impact:** Educating users about prompt hijacking is crucial, but it can be challenging to reach a broad audience effectively.

7. Detection and Mitigation

•**Challenge:** Developing effective methods to detect and mitigate prompt hijacking is a complex task. This involves creating robust input validation systems that can filter out harmful prompts without affecting the model's performance.

•**Impact:** The challenge lies in balancing safety with usability, ensuring that legitimate prompts are not mistakenly flagged while harmful ones are effectively identified.

ASSIGNMENT – UNIT V

1. What is In-Context Learning (ICL)? Discuss its characteristics and how it differs from traditional machine learning methods.
2. Define prompt hijacking. How does it occur in language models, and what ethical challenges does it present? Suggest strategies to mitigate these risks.
3. What are the various techniques for effective prompt engineering? Provide examples to illustrate how these techniques can enhance the performance of language models.
4. Explain image prompting. How can it be used alongside text prompts, and what are the potential benefits and challenges associated with this technique?
5. Compare and contrast in-context prompting with traditional prompting methods. What are the advantages and limitations of each approach? In what contexts is each method most effective?

PART A- UNIT-V

1.What is prompt engineering?

Answer: Prompt engineering is the process of crafting text prompts to help large language models (LLMs) generate more accurate and creative outputs.

2.Define a prompt in the context of AI models.

Answer: A prompt is an input instruction or cue that shapes the model's response, guiding it to generate relevant and informative output.

3.What are natural language prompts?

Answer: Natural language prompts emulate human-like instructions, providing guidance in natural language, making interaction with models more intuitive.

4.Explain the purpose of system prompts.

Answer: System prompts are predefined instructions or templates that guide the model's output, specifying the desired output format or behavior.

5.What are conditional prompts?

Answer: Conditional prompts involve conditioning the model on specific contexts or constraints using statements like "If X, then Y."

6.How does prompt engineering influence an LLM's output?

Answer: By carefully choosing words and phrases in a prompt, prompt engineering influences how an LLM interprets a task and the results it produces.

PART A- UNIT-V

7.What is the first step in effective prompt engineering?

Answer: The first step is to understand the task clearly and define what output is expected from the LLM.

8.Why is clear and concise language important in prompt engineering?

Answer: Clear and concise language helps ensure that the LLM understands the prompt without ambiguity, avoiding confusion.

9.How can specificity enhance prompt effectiveness?

Answer: Being specific in prompts increases the likelihood of generating relevant and informative responses from the LLM.

10.What role do examples play in prompt engineering?

Answer: Providing examples helps the LLM understand the kind of output being sought, improving the quality of its responses.

11.What are some common challenges in prompt engineering?

Answer: Common challenges include ambiguity in prompts, unexpected model behavior, and difficulty in achieving the desired output format.

12.Explain the concept of prompt hijacking.

Answer: Prompt hijacking refers to the manipulation of prompts to produce unintended or harmful outputs from an AI model.

13. What techniques can be used to mitigate prompt hijacking?

Answer: Techniques include refining prompt language, implementing user input validation, and employing robust monitoring systems.

PART A- UNIT-V

14.What is the significance of using diverse types of prompts?

Answer: Using diverse types of prompts helps in maximizing the model's performance by catering to different contexts and requirements.

15.How can image prompting enhance AI interactions?

Answer: Image prompting allows models to understand and generate responses based on visual cues, enriching the interaction experience.

16.What is the difference between natural language prompts and conditional prompts?

Answer: Natural language prompts provide straightforward instructions, while conditional prompts set specific conditions for model responses.

17. Why is it important to avoid jargon in prompts?

Answer: Avoiding jargon ensures that the prompt is accessible and easily understood by the LLM, facilitating better responses.

18. Describe how the use of structured prompts can guide model behavior.

Answer: Structured prompts provide clear formats and guidelines, helping the model to follow expected patterns in generating output.

19. What are some applications of prompt engineering in real-world scenarios?

Answer: Applications include chatbots, content generation, data analysis, and creative writing, where clear guidance is essential.

PART A- UNIT-V

and poor performance often encountered when training directly at high resolutions.

18. What are some common applications of Paired Style Transfer?

Paired Style Transfer is commonly used in applications such as image-to-image translation, for example, converting sketches to photos, black-and-white images to color, or day-to-night transformations in photography.

19. How are GANs used in style transfer tasks?

GANs are used in style transfer to learn the mapping between two domains, such as transferring artistic styles to photographs or converting images from one domain (e.g., summer) to another (e.g., winter).

20. What are the main modes of operation in deepfakes?

The two primary modes of operation in deepfakes are replacement, where one person's face is swapped with another's, and re-enactment, where a person's expressions or movements are modified to produce new behaviors in a video.

PART B- UNIT V

- 1.What is prompt engineering, and why is it important in the context of large language models?**
- 2. Explain the difference between natural language prompts and conditional prompts. Provide examples for each.**
- 3.Discuss the role of specificity in prompt design. How can being specific enhance the effectiveness of a prompt?**
- 4.What challenges might arise in prompt engineering, and what strategies can be employed to overcome these challenges?**
- 5.Describe the concept of image prompting and how it can enhance AI interactions. What are some potential applications?**
- 6.How can understanding user intent improve the design of prompts in AI systems? Provide an example to illustrate your point.**

SUPPORTIVE ONLINE COURSES – UNIT V

- <https://www.coursera.org/learn/prompt-engineering>
- <https://www.udemy.com/topic/prompt-engineering/>
- <https://www.simplilearn.com/prompt-engineering-free-course-skillup>
- <https://www.datacamp.com/>



Real Time Applications in Day to Day life and to Industry

1. Virtual Assistants (In-Context Learning & Prompting in Day-to-Day Life)

Virtual assistants like **Google Assistant**, **Alexa**, or **Siri** use **In-Context Learning** to provide more accurate answers or perform tasks based on previous conversations or queries. This happens in real time, meaning as you continue interacting, they understand your preferences better.

For example, if you frequently ask your virtual assistant to set an alarm for 7 AM, it may prompt you with "Would you like me to set your usual 7 AM alarm?" without needing you to specify it every time.



2. Chatbots for Customer Support (In-Context Prompting in Industry)

In the **industry**, chatbots deployed for customer support can apply **In-Context Prompting** techniques to understand the flow of a conversation and respond in real time. These bots are used across **e-commerce** platforms, **banks**, and **telecommunications** companies. When a customer starts a conversation, the bot remembers previous interactions and suggests solutions accordingly.

For example, if a customer recently asked about their order status, the bot might immediately suggest, "Do you want to check your current order status?" when they return.



PREScribed TEXT BOOKS AND REFERENCE BOOKS

TEXT BOOKS:

1. "GANs in Action: Deep Learning with Generative Adversarial Networks" by Jakub Langr and Vladimir Bok
2. "Deep Learning with Python" by François Chollet
3. "Deep Learning with TensorFlow and Keras" by Amita Kapoor

REFERENCES:

1. Goodfellow, Ian, et al. ["Generative Adversarial Nets"](#), 2014.
2. Karras, Tero, et al. ["Progressive Growing of GANs for Improved Quality, Stability, and Variation"](#), 2018.
3. Isola, Phillip, et al. ["Image-to-Image Translation with Conditional Adversarial Networks"](#), 2017.

MINI PROJECT SUGGESTIONS

- ✓ Team 1: Develop a Basic In-Context Learning Model for Text Prediction
- ✓ Team 2: Create a Contextual AI Chatbot with Adaptive Response
- ✓ Team 3: Explore Advanced In-Context Learning Techniques in Language Models
- ✓ Team 4: Develop a GAN Model for Image Prompting and Generation
- ✓ Team 5: Investigate Prompt Hijacking in NLP Models
- ✓ Team 6: Address Challenges in In-Context Learning with Continual Learning

THANK YOU

Disclaimer:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

R.M.K GROUP OF ENGINEERING INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS

R.M.K GROUP OF INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS



Please read this disclaimer before proceeding:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

Digital Course Material
22AI005
Introduction to Generative AI



Department : CSE(CYBER SECURITY)

Batch/Year : 2022-2026/IV

Created by : Ms. G.K. MONICA

AP/CSE(CS)

Date : 07-07-2025

TABLE OF CONTENTS

S.NO	CONTENTS	SLIDE NO.
1	COURSE OBJECTIVES	6
2	PRE REQUISITES (COURSE NAMES WITH CODE)	7
3	SYLLABUS (WITH SUBJECT CODE, NAME, LTPC DETAILS)	8
4	COURSE OUTCOMES (6)	9
5	CO- PO/PSO MAPPING	10
6	LECTURE PLAN –UNIT 1	12
7	ACTIVITY BASED LEARNING –UNIT 1	14
8	TEST YOURSELF	15
9	LECTURE NOTES	16
10	ASSIGNMENT	50
11	PART A Q s (WITH K LEVEL AND CO)	51
12	PART B Q s (WITH K LEVEL AND CO)	53
13	SUPPORTIVE ONLINE CERTIFICATION COURSES	54
14	REAL TIME APPLICATIONS IN DAY TO DAY LIFE AND TO INDUSTRY	55
15	PRESCRIBED TEXT BOOKS & REFERENCE BOOKS	59
16	MINI PROJECT SUGGESTIONS	60
17	TOUGH QUESTIONS SIMILAR TO GATE	61

COURSE OBJECTIVES


The Course will enable learners to:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.



PRE REQUISITES

✿ PRE-REQUISITE CHART



OBJECTIVES:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.

UNIT I INTRODUCTION**9**

Generative Models – Image transformation – Challenges - Deep Neural Networks – Perceptron – back propagation – CNN – RNN – Optimizer.

UNIT II IMAGE GENERATION**9**

Creating encodings of images – variational objective – Inverse Autoregressive flow – Importing CIFAR – Creating the network from TensorFlow 2.

UNIT III GENERATIVE ADVERSARIAL NETWORKS**9**

Generative Adversarial Networks – Vanilla GAN – Improved GANs – Progressive GAN – Challenges – Paired style transfer – Unpaired style transfer – Deepfakes – Modes of operation – key feature set – High level flow – Replacement – Re-enactment.

UNIT IV LARGE LANGUAGE MODELS**9**

Overview of LLMs - Transformers – GPT – Types of LLMs – Key concepts – other Transformers – T5 – Generative Pre-Training Models – Multi-modal Models – DALL.E 2

UNIT V PROMPT ENGINEERING**9**

Basics – In-Context Learning – In-Context Prompting – Techniques – Image Prompting – Prompt Hijacking – Challenges.

TOTAL: 45 PERIODS

COURSE OUTCOME

Course Code	Course Outcome Statement	Cognitive/ Affective Level of the Course Outcome	Expected Level of Attainment
Course Outcome Statements in Cognitive Domain			
C305.1	Elaborate the basic concepts of Generative AI	Understand K2	60%
C305.2	Build Generative AI systems to generate images	Analyse K4	60%
C305.3	Apply the concepts used in Generative AI Models	Apply K3	60%
C305.4	Use various Generative AI models.	Apply K3	60%
C305.5	Compare and use the various Large Language Models	Analyse K4	60%
C305.6	Analyze the basics of Prompt Engineering.	Apply K3	60%
Course Outcome Statements in Affective domain			
C305.7	Attend the classes regularly	Respond (A2)	95%
C305.8	Submit the Assignments regularly.	Respond (A2)	95%
C305.9	Participation in Seminar/Quiz/ Group Discussion/ Collaborative learning and content beyond syllabus	Valuing (A3)	95%

CO-PO/PSO MAPPING

Correlation Matrix of the Course Outcomes to Programme Outcomes and Programme Specific Outcomes Including Course Enrichment Activities

Course Outcomes (Cos)		Programme Outcomes (POs), Programme Specific Outcomes (PSOs)														
		PO 1	PO 2	PO 3	PO 4	PO 5	PO 6	PO 7	PO 8	PO 9	PO 10	PO 11	PO 12	PSO 1	PSO 2	PSO 3
		K 3	K4	K 5	K5	K3 /K 5	A2	A3	A 3	A3	A3	A3	A2	K3	K3	K3
C305.1	K 3	3	2	1	1	3						3		3	3	3
C305.2	K 4	3	3	2	2	3						3		3	3	3
C305.3	K 2	2	1									2		3	3	3
C305.4	K 3	3	2	1	1	3						3		3	3	3
C305.5	K 4	3	3	2	2	3						3		3	3	3
C305.6	K 3	3	2	1	1	3						3		2	2	2
C305.7	A 2												3			
C305.8	A 2								2	2	2		3			
C305.9	A 3						3	3		3	3		3			
C305		3	3	2	2	3	1	1	1	3	3	3	3	3	3	3



UNIT I

INTRODUCTION

LECTURE PLAN – UNIT I

	UNIT I INTRODUCTION						
Sl. No	TOPIC	NO OF PERIODS	PROPOSED LECTURE	ACTUAL LECTURE	PERTAINING CO(s)	TAXONOMY LEVEL	MODE OF DELIVERY
			PERIOD	PERIOD			
1	Generative Models	1			CO1	K2	PPT
2	Image transformation	1			CO1	K2	PPT
3	Challenges	1			CO1	K2	PPT
4	Deep Neural Networks	1			CO1	K3	PPT
5	Perceptron	1			CO1	K3	PPT
6	Back propagation	1			CO1	K2	PPT
7	CNN	1			CO1	K2	PPT
8	RNN	1			CO1	K2	PPT
9	Optimizer.	1			CO1	K2	PPT

LECTURE PLAN – UNIT I

❁ ASSESSMENT COMPONENTS

- ❁ AC 1. Unit Test
- ❁ AC 2. Assignment
- ❁ AC 3. Course Seminar
- ❁ AC 4. Course Quiz
- ❁ AC 5. Case Study
- ❁ AC 6. Record Work
- ❁ AC 7. Lab / Mini Project
- ❁ AC 8. Lab Model Exam
- ❁ AC 9. Project Review

MODE OF DELEIVERY

- MD 1. Oral presentation
- MD 2. Tutorial
- MD 3. Seminar
- MD 4 Hands On
- MD 5. Videos
- MD 6. Field Visit



R.M.K.
GROUP OF
INSTITUTIONS

ACTIVITY BASED LEARNING – UNIT I

QUIZ- LINK

Unit I:

Generative Models

<https://custom.typingmind.com/tools/gen-ai-quiz>

Deep Neural Networks

<https://www.linkedin.com/advice/0/what-best-online-quizzes-testing-your-neural-network>

CNN

<https://quizizz.com/admin/quiz/5f2544e17c727f001b6b97c4/convolutional-neural-network>

RNN

<https://github.com/abhishektripathi24/Deep-Learning-Specialization-Coursera/blob/master/5.%20Sequence-Models/week1/Quiz%20-%20Recurrent%20Neural%20Networks.pdf>

Perceptrons and Backpropagation

<https://webcms3.cse.unsw.edu.au/COMP9444/19T3/resources/34909>

TEST YOURSELF

1. What are some of the potential benefits of Generative AI?

- A. Generative AI can be used to create new and innovative products and services.
- B. Generative AI can be used to improve the quality of life for people with disabilities.
- C. Generative AI can be used to solve complex problems that are currently beyond the reach of human intelligence.
- D. All of the above

2. Which of the following is a subset of machine learning?

- A. SciPy
- B. NumPy
- C. Deep learning
- D. All of the above

3. A perceptron is a _____

- A. Backtracking algorithm
- B. Backpropagation algorithm
- C. Feed-forward neural network
- D. Feed Forward-backward algorithm

4. What is a Convolutional Neural Network (CNN)?

- A. A type of recurrent neural network
- B. An unsupervised learning algorithm
- C. A deep reinforcement learning model
- D. A specialized neural network for image processing

5. What is the primary purpose of a Recurrent Neural Network (RNN)?

- A) Image classification
- B) Text generation
- C) Reinforcement learning
- D) Object detection

UNIT I

INTRODUCTION

1.1 Generative Models

A generative model is a type of machine learning model that aims to learn the underlying patterns or distributions of data in order to generate new, similar data. In essence, it's like teaching a computer to dream up its own data based on what it has seen before. The significance of this model lies in its ability to create, which has vast implications in various fields, from art to science.

What is a Generative Model?

Generative models use machine learning to discover patterns in data & generate new data. Learn about their significance & applications in AI.

Generative Models Explained

Generative models are a cornerstone in the world of artificial intelligence (AI). Their primary function is to understand and capture the underlying patterns or distributions from a given set of data. Once these patterns are learned, the model can then generate new data that shares similar characteristics with the original dataset.

Imagine you're teaching a child to draw animals. After showing them several pictures of different animals, the child begins to understand the general features of each animal. Given some time, the child might draw an animal they've never seen before, combining features they've learned. This is analogous to how a generative model operates: it learns from the data it's exposed to and then creates something new based on that knowledge.

The distinction between generative and discriminative models is fundamental in machine learning:

Generative models: These models focus on understanding how the data is generated. They aim to learn the distribution of the data itself. For instance, if we're looking at pictures of cats and dogs, a generative model would try to understand what makes a cat look like a cat and a dog look like a dog. It would then be able to generate new images that resemble either cats or dogs.

Discriminative models: These models, on the other hand, focus on distinguishing between different types of data. They don't necessarily learn or understand how the data is generated; instead, they learn the boundaries that separate one class of data from another. Using the same example of cats and dogs, a discriminative model would learn to tell the difference between the two, but it wouldn't necessarily be able to generate a new image of a cat or dog on its own.

In the realm of AI, generative models play a pivotal role in tasks that require the creation of new content. This could be in the form of synthesizing realistic human faces, composing music, or even generating textual content. Their ability to "dream up" new data makes them invaluable in scenarios where original content is needed, or where the augmentation of existing datasets is beneficial.

In essence, while discriminative models excel at classification tasks, generative models shine in their ability to create. This creative prowess, combined with their deep understanding of data distributions, positions generative models as a powerful tool in the AI toolkit.

Types of Generative Models

Generative models come in various forms, each with its unique approach to understanding and generating data. Here's a more comprehensive list of some of the most prominent types:

Bayesian networks. These are graphical models that represent the probabilistic relationships among a set of variables. They're particularly useful in scenarios where understanding causal relationships is crucial. For example, in medical diagnosis, a Bayesian network might help determine the likelihood of a disease given a set of symptoms.

Diffusion models. These models describe how things spread or evolve over time. They're often used in scenarios like understanding how a rumor spreads in a network or predicting the spread of a virus in a population.

Generative Adversarial Networks (GANs). GANs consist of two neural networks, the generator and the discriminator, that are trained together. The generator tries to produce data, while the discriminator attempts to distinguish between real and generated data. Over time, the generator becomes so good that the discriminator can't tell the difference. GANs are popular in image generation tasks, such as creating realistic human faces or artworks.

Variational Autoencoders (VAEs). VAEs are a type of autoencoder that produces a compressed representation of input data, then decodes it to generate new data. They're often used in tasks like image denoising or generating new images that share characteristics with the input data.

Restricted Boltzmann Machines (RBMs). RBMs are neural networks with two layers that can learn a probability distribution over its set of inputs. They've been used in recommendation systems, like suggesting movies on streaming platforms based on user preferences.

Types of Generative Models

Pixel Recurrent Neural Networks (PixelRNNs). These models generate images pixel by pixel, using the context of previous pixels to predict the next one. They're particularly useful in tasks where the sequential generation of data is crucial, like drawing an image line by line.

Markov chains. These are models that predict future states based solely on the current state, without considering the states that preceded it. They're often used in text generation, where the next word in a sentence is predicted based on the current word.

Normalizing flows. These are a series of invertible transformations applied to simple probability distributions to produce more complex distributions. They're useful in tasks where understanding the transformation of data is crucial, like in financial modeling.

Real-World Use Cases of Generative Models

Generative models have penetrated mainstream consumption, revolutionizing the way we interact with technology and experience content, for example:

- **Art creation.** Artists and musicians are using generative models to create new pieces of art or compositions, based on styles they feed into the model. For example, [Midjourney](#) is a very popular tool that is used to generate artwork.
- **Drug discovery.** Scientists can use generative models to predict molecular structures for new potential drugs.
- **Content creation.** Website owners leverage generative models to speed up the content creation process. For example, [Hubspot's AI content writer](#) helps marketers generate blog posts, landing page copy and social media posts.
- **Video games.** Game designers use generative models to create diverse and unpredictable game environments or characters.

What are the Benefits of Generative Models?

Generative models, with their unique ability to create and innovate, offer a plethora of advantages that extend beyond mere data generation. Here's a deeper dive into the myriad benefits they bring to the table:

- **Data augmentation.** In domains where data is scarce or expensive to obtain, generative models can produce additional data to supplement the original set. For instance, in medical imaging, where obtaining large datasets can be challenging, these models can generate more images to aid in better training of diagnostic tools.
- **Anomaly detection.** By gaining a deep understanding of what constitutes "normal" data, generative models can efficiently identify anomalies or outliers. This is particularly useful in sectors like finance, where spotting fraudulent transactions quickly is paramount.
- **Flexibility.** Generative models are versatile and can be employed in a range of learning scenarios, including [unsupervised](#), semi-supervised, and [supervised learning](#). This adaptability makes them suitable for a wide array of tasks.
- **Personalization.** These models can be tailored to generate content based on specific user preferences or inputs. For example, in the entertainment industry, generative models can create personalized music playlists or movie recommendations, enhancing user experience.
- **Innovation in design.** In fields like architecture or product design, generative models can propose novel designs or structures, pushing the boundaries of creativity and innovation.
- **Cost efficiency.** By automating the creation of content or solutions, generative models can reduce the costs associated with manual production or research, leading to more efficient processes in industries like manufacturing or entertainment.

What are the Limitations of Generative Models?

While generative models are undeniably powerful and transformative, they are not without their challenges. Here's an exploration of some of the constraints and challenges associated with these models:

- **Training complexity.** Generative models, especially sophisticated ones like GANs, require significant computational resources and time. Training them demands powerful hardware and can be resource-intensive.
- **Quality control.** While they can produce vast amounts of data, ensuring the quality and realism of the generated content can be challenging. For instance, a model might generate an image that looks realistic at first glance but has subtle anomalies upon closer inspection.
- **Overfitting.** There's a risk that generative models can become too attuned to the training data, producing outputs that lack diversity or are too closely tied to the input they've seen.
- **Lack of interpretability.** Many generative models, particularly deep learning-based ones, are often seen as "black boxes." This means it can be challenging to understand how they make decisions or why they produce specific outputs, which can be a concern in critical applications like healthcare.
- **Ethical concerns.** The ability of generative models to produce realistic content raises ethical issues, especially in the creation of deep fakes or counterfeit content. Ensuring responsible use is paramount to prevent misuse or deception.
- **Data dependency.** The quality of the generated output is heavily dependent on the quality of the training data. If the training data is biased or unrepresentative, the model's outputs will reflect those biases.
- **Mode collapse.** Particularly in GANs, there's a phenomenon called mode collapse where the generator produces limited varieties of samples, reducing the diversity of the generated outputs.

Image transformation

Image transformation in the context of generating AI refers to the ability of artificial intelligence models, particularly generative models like Generative Adversarial Networks (GANs) or Variational Autoencoders (VAEs), to alter or enhance images in various ways. Here's a detailed look at how image transformation is approached in generating AI:

Techniques and Approaches

1. Generative Adversarial Networks (GANs):

Overview: GANs consist of two neural networks — a generator and a discriminator — that are trained adversarial to generate realistic images.

Image Transformation: GANs can transform images in several ways:

- 1. Style Transfer:** Altering the artistic style of an image while preserving its content.
- 2. Image-to-Image Translation:** Converting images from one domain to another (e.g., day to night, sketches to photographs).
- 3. Super-Resolution:** Increasing the resolution and quality of low-resolution images.
- 4. Face Aging and Rejuvenation:** Predicting how faces might look older or younger based on learned features.
- 5. Inpainting:** Filling in missing parts of an image based on surrounding context.

2. Variational Autoencoders (VAEs):

Overview: VAEs consist of an encoder network that maps input images into a latent space and a decoder network that reconstructs images from this latent space.

Image Transformation: VAEs can transform images by modifying their latent representations:

- 1. Image Generation:** Generating new images from learned latent representations.
- 2. Image Editing:** Interpolating between latent representations to edit specific attributes (e.g., changing colors, styles).
- 3. Cross-Domain Transformation:** Learning shared latent spaces to translate images between different domains (e.g., maps to satellite images).

Challenges and Considerations

1. Quality and Realism:

1. Ensuring that generated images are realistic and of high quality, without artifacts or distortions.
2. Maintaining fidelity to the input while applying transformations.

2. Semantic Understanding:

1. Preserving semantic meaning and context during transformations (e.g., not distorting essential features or objects).
2. Handling multi-modal outputs where multiple plausible transformations exist for a single input.

3. Control and Interactivity:

1. Providing user control over specific aspects of transformations (e.g., style, attributes).
2. Enabling interactive manipulation of images in real-time applications.

4. Generalization and Robustness:

1. Ensuring that models generalize well to unseen data and diverse input conditions.
2. Addressing biases and ensuring fairness in transformations, particularly in sensitive applications like face manipulation.

5. Computational Efficiency:

1. Optimizing models for fast inference and real-time applications.
2. Managing computational resources required for training and deployment, especially for large-scale datasets.

Applications

Image transformation in generating AI has diverse applications across various fields:


- **Creative Arts:** Enabling artists to explore new styles and create innovative artwork.
- **Entertainment:** Enhancing visual effects in movies and games.
- **Healthcare:** Generating synthetic medical images for training and diagnostics.
- **Design and Fashion:** Facilitating virtual try-ons and customization.
- **Education and Training:** Creating realistic simulations and educational materials.

Deep Neural Networks

A deep neural network (DNN) is an ANN with multiple hidden layers between the input and output layers. Similar to shallow ANNs, DNNs can model complex non-linear relationships.

The main purpose of a neural network is to receive a set of inputs, perform progressively complex calculations on them, and give output to solve real world problems like classification. We restrict ourselves to feed forward neural networks.

We have an input, an output, and a flow of sequential data in a deep network.



Neural networks are widely used in supervised learning and reinforcement learning problems. These networks are based on a set of layers connected to each other.

In deep learning, the number of hidden layers, mostly non-linear, can be large; say about 1000 layers.

DL models produce much better results than normal ML networks.


We mostly use the gradient descent method for optimizing the network and minimizing the loss function.

We can use the Imagenet, a repository of millions of digital images to classify a dataset into categories like cats and dogs. DL nets are increasingly used for dynamic images apart from static ones and for time series and text analysis.

Training the data sets forms an important part of Deep Learning models. In addition, Backpropagation is the main algorithm in training DL models.

DL deals with training large neural networks with complex input output transformations.


One example of DL is the mapping of a photo to the name of the person(s) in photo as they do on social networks and describing a picture with a phrase is another recent application of DL.



Neural networks are functions that have inputs like $x_1, x_2, x_3, \dots$ that are transformed to outputs like z_1, z_2, z_3 and so on in two (shallow networks) or several intermediate operations also called layers (deep networks).

The weights and biases change from layer to layer. 'w' and 'v' are the weights or synapses of layers of the neural networks.

The best use case of deep learning is the supervised learning problem. Here, we have a large set of data inputs with a desired set of outputs.



Here we apply back propagation algorithm to get correct output prediction.

The most basic data set of deep learning is the MNIST, a dataset of handwritten digits.

We can train deep a Convolutional Neural Network with Keras to classify images of handwritten digits from this dataset.

The firing or activation of a neural net classifier produces a score. For example, to classify patients as sick and healthy, we consider parameters such as height, weight and body temperature, blood pressure etc.

A high score means patient is sick and a low score means he is healthy.

Each node in output and hidden layers has its own classifiers. The input layer takes inputs and passes on its scores to the next hidden layer for further activation and this goes on till the output is reached.

This progress from input to output from left to right in the forward direction is called forward propagation.

Credit assignment path (CAP) in a neural network is the series of transformations starting from the input to the output. CAPs elaborate probable causal connections between the input and the output.

CAP depth for a given feed forward neural network or the CAP depth is the number of hidden layers plus one as the output layer is included. For recurrent neural networks, where a signal may propagate through a layer several times, the CAP depth can be potentially limitless.

Perceptron

A single-layer feedforward neural network was introduced in the late 1950s by Frank Rosenblatt. It was the starting phase of Deep Learning and Artificial neural networks. During that time for prediction, Statistical machine learning, or Traditional code Programming is used. Perceptron is one of the first and most straightforward models of artificial neural networks. Despite being a straightforward model, the perceptron has been proven to be successful in solving specific categorization issues.

What is Perceptron?

Perceptron is one of the simplest Artificial neural network architectures. It was introduced by Frank Rosenblatt in 1957s. It is the simplest type of feedforward neural network, consisting of a single layer of input nodes that are fully connected to a layer of output nodes. It can learn the linearly separable patterns. it uses slightly different types of artificial neurons known as threshold logic units (TLU). it was first introduced by McCulloch and Walter Pitts in the 1940s.

Types of Perceptron

Single-Layer Perceptron: This type of perceptron is limited to learning linearly separable patterns. effective for tasks where the data can be divided into distinct categories through a straight line.

Multilayer Perceptron: Multilayer perceptrons possess enhanced processing capabilities as they consist of two or more layers, adept at handling more complex patterns and relationships within the data.

Basic Components of Perceptron

A perceptron, the basic unit of a neural network, comprises essential components that collaborate in information processing.

Input Features: The perceptron takes multiple input features, each input feature represents a characteristic or attribute of the input data.

Weights: Each input feature is associated with a weight, determining the significance of each input feature in influencing the perceptron's output. During training, these weights are adjusted to learn the optimal values.

Summation Function: The perceptron calculates the weighted sum of its inputs using the summation function. The summation function combines the inputs with their respective weights to produce a weighted sum.

Activation Function: The weighted sum is then passed through an activation function. Perceptron uses Heaviside step function functions. which take the summed values as input and compare with the threshold and provide the output as 0 or 1.

Output: The final output of the perceptron, is determined by the activation function's result. For example, in binary classification problems, the output might represent a predicted class (0 or 1).

Bias: A bias term is often included in the perceptron model. The bias allows the model to make adjustments that are independent of the input. It is an additional parameter that is learned during training.

Learning Algorithm (Weight Update Rule): During training, the perceptron learns by adjusting its weights and bias based on a learning algorithm. A common approach is the perceptron learning algorithm, which updates weights based on the difference between the predicted output and the true output.

These components work together to enable a perceptron to learn and make predictions. While a single perceptron can perform binary classification, more complex tasks require the use of multiple perceptrons organized into layers, forming a neural network.

How does Perceptron work?


A weight is assigned to each input **node** of a perceptron, indicating the significance of that input to the output. The perceptron's output is a weighted sum of the inputs that have been run through an activation function to decide whether or not the perceptron will fire. it computes the weighted sum of its inputs as:

$$Z = w_1X_1 + w_2X_2 + \dots + w_nX_n = X^TW$$

The step function compares this weighted sum to the threshold, which outputs 1 if the input is larger than a threshold value and 0 otherwise, is the activation function that perceptrons utilize the most frequently. The most common step function used in perceptron is the Heaviside step function:

$$h(z) = \begin{cases} 0 & \text{if } z < \text{Threshold} \\ 1 & \text{if } z \geq \text{Threshold} \end{cases}$$

A perceptron has a single layer of threshold logic units with each TLU connected to all inputs.



When all the neurons in a layer are connected to every neuron of the previous layer, it is known as a fully connected layer or dense layer.

The output of the fully connected layer can be:

$$f_{\{W,b\}}(X) = h(XW + b)$$

where X is the input W is the weight for each inputs neurons and b is the bias and h is the step function.

During training, The perceptron's weights are adjusted to minimize the difference between the predicted output and the actual output. Usually, supervised learning algorithms like the delta rule or the perceptron learning rule are used for this.

$$w_{\{i,j\}} = w_{\{i,j\}} + \eta (y_j - \hat{y}_j)x_i$$

Here $w_{i,j}$ is the weight between the i th input and j th output neuron, x_i is the i th input value, and y_j and $\hat{y}_j$ is the j th actual and predicted value is η the learning rate.

Backpropagation


Machine learning models learn from data and make predictions. One of the fundamental concepts behind training these models is backpropagation. In this article, we will explore what backpropagation is, why it is crucial in machine learning, and how it works.

What is backpropagation?

In machine learning, backpropagation is an effective algorithm used to train artificial neural networks, especially in feed-forward neural networks.

Backpropagation is an iterative algorithm, that helps to minimize the cost function by determining which weights and biases should be adjusted. During every epoch, the model learns by adapting the weights and biases to minimize the loss by moving down toward the gradient of the error. Thus, it involves the two most popular optimization algorithms, such as gradient descent or stochastic gradient descent.

Computing the gradient in the backpropagation algorithm helps to minimize the cost function and it can be implemented by using the mathematical rule called chain rule from calculus to navigate through complex layers of the neural network.



Advantages of Using the Backpropagation Algorithm in Neural Networks

Backpropagation, a fundamental algorithm in training neural networks, offers several advantages that make it a preferred choice for many machine learning tasks. Here, we discuss some key advantages of using the backpropagation algorithm:

Ease of Implementation: Backpropagation does not require prior knowledge of neural networks, making it accessible to beginners. Its straightforward nature simplifies the programming process, as it primarily involves adjusting weights based on error derivatives.

Simplicity and Flexibility: The algorithm's simplicity allows it to be applied to a wide range of problems and network architectures. Its flexibility makes it suitable for various scenarios, from simple feedforward networks to complex recurrent or convolutional neural networks.

Efficiency: Backpropagation accelerates the learning process by directly updating weights based on the calculated error derivatives. This efficiency is particularly advantageous in training deep neural networks, where learning features of a function can be time-consuming.

Generalization: Backpropagation enables neural networks to generalize well to unseen data by iteratively adjusting weights during training. This generalization ability is crucial for developing models that can make accurate predictions on new, unseen examples.

Scalability: Backpropagation scales well with the size of the dataset and the complexity of the network. This scalability makes it suitable for large-scale machine learning tasks, where training data and network size are significant factors.

In conclusion, the backpropagation algorithm offers several advantages that contribute to its widespread use in training neural networks. Its ease of implementation, simplicity, efficiency, generalization ability, and scalability make it a valuable tool for developing and training neural network models for various machine learning applications.

Working of Backpropagation Algorithm

The Backpropagation algorithm works by two different passes, they are:


- **Forward pass**
- **Backward pass**

How does Forward pass work?

In forward pass, initially the input is fed into the input layer. Since the inputs are raw data, they can be used for training our neural network.

The inputs and their corresponding weights are passed to the hidden layer. The hidden layer performs the computation on the data it receives. If there are two hidden layers in the neural network, for instance, consider the illustration fig(a), h_1 and h_2 are the two hidden layers, and the output of h_1 can be used as an input of h_2 . Before applying it to the activation function, the bias is added.

To the weighted sum of inputs, the activation function is applied in the hidden layer to each of its neurons. One such activation function that is commonly used is ReLU can also be used, which is responsible for returning the input if it is positive otherwise it returns zero. By doing this so, it introduces the non-linearity to our model, which enables the network to learn the complex relationships in the data. And finally, the weighted outputs from the last hidden layer are fed into the output layer to compute the final prediction, this layer can also use the activation function called the softmax function which is responsible for converting the weighted outputs into probabilities for each class.



How does backward pass work?

In the backward pass process shows, the error is transmitted back to the network which helps the network, to improve its performance by learning and adjusting the internal weights.

To find the error generated through the process of forward pass, we can use one of the most commonly used methods called mean squared error which calculates the difference between the predicted output and desired output. The formula for mean squared error is:

$$\text{Meansquarederror} = (\text{predictedoutput} - \text{actualoutput})^2$$

Once we have done the calculation at the output layer, we then propagate the error backward through the network, layer by layer.

The key calculation during the backward pass is determining the gradients for each weight and bias in the network. This gradient is responsible for telling us how much each weight/bias should be adjusted to minimize the error in the next forward pass. The chain rule is used iteratively to calculate this gradient efficiently.

In addition to gradient calculation, the activation function also plays a crucial role in backpropagation, it works by calculating the gradients with the help of the derivative of the activation function.

Convolution Neural Network

A Convolutional Neural Network (CNN) is a type of Deep Learning neural network architecture commonly used in Computer Vision. Computer vision is a field of Artificial Intelligence that enables a computer to understand and interpret the image or visual data.

When it comes to Machine Learning, Artificial Neural Networks perform really well. Neural Networks are used in various datasets like images, audio, and text. Different types of Neural Networks are used for different purposes, for example for predicting the sequence of words we use Recurrent Neural Networks more precisely an LSTM, similarly for image classification we use Convolution Neural networks. In this blog, we are going to build a basic building block for CNN.

In a regular Neural Network there are three types of layers:

1.Input Layers: It's the layer in which we give input to our model. The number of neurons in this layer is equal to the total number of features in our data (number of pixels in the case of an image).

2.Hidden Layer: The input from the Input layer is then fed into the hidden layer. There can be many hidden layers depending on our model and data size. Each hidden layer can have different numbers of neurons which are generally greater than the number of features. The output from each layer is computed by matrix multiplication of the output of the previous layer with learnable weights of that layer and then by the addition of learnable biases followed by activation function which makes the network nonlinear.

3.Output Layer: The output from the hidden layer is then fed into a logistic function like sigmoid or softmax which converts the output of each class into the probability score of each class.


The data is fed into the model and output from each layer is obtained from the above step is called feedforward, we then calculate the error using an error function, some common error functions are cross-entropy, square loss error, etc. The error function measures how well the network is performing. After that, we backpropagate into the model by calculating the derivatives. This step is called Backpropagation which basically is used to minimize the loss.

Convolution Neural Network

Convolutional Neural Network (CNN) is the extended version of artificial neural networks (ANN) which is predominantly used to extract the feature from the grid-like matrix dataset. For example visual datasets like images or videos where data patterns play an extensive role.

CNN architecture


Convolutional Neural Network consists of multiple layers like the input layer, Convolutional layer, Pooling layer, and fully connected layers.




The Convolutional layer applies filters to the input image to extract features, the Pooling layer downsamples the image to reduce computation, and the fully connected layer makes the final prediction. The network learns the optimal filters through backpropagation and gradient descent.

How Convolutional Layers works

Convolution Neural Networks or convnets are neural networks that share their parameters. Imagine you have an image. It can be represented as a cuboid having its length, width (dimension of the image), and height (i.e the channel as images generally have red, green, and blue channels).



Now imagine taking a small patch of this image and running a small neural network, called a filter or kernel on it, with say, K outputs and representing them vertically. Now slide that neural network across the whole image, as a result, we will get another image with different widths, heights, and depths. Instead of just R, G, and B channels now we have more channels but lesser width and height. This operation is called Convolution. If the patch size is the same as that of the image it will be a regular net



Now let's talk about a bit of mathematics that is involved in the whole convolution process.

- Convolution layers consist of a set of learnable filters (or kernels) having small widths and heights and the same depth as that of input volume (3 if the input layer is image input).
- For example, if we have to run convolution on an image with dimensions $34 \times 34 \times 3$. The possible size of filters can be $a \times a \times 3$, where 'a' can be anything like 3, 5, or 7 but smaller as compared to the image dimension.
- During the forward pass, we slide each filter across the whole input volume step by step where each step is called stride (which can have a value of 2, 3, or even 4 for high-dimensional images) and compute the dot product between the kernel weights and patch from input volume.
- As we slide our filters we'll get a 2-D output for each filter and we'll stack them together as a result, we'll get output volume having a depth equal to the number of filters. The network will learn all the filters.

Layers used to build ConvNets

A complete Convolution Neural Networks architecture is also known as convnets. A convnets is a sequence of layers, and every layer transforms one volume to another through a differentiable function.


Types of layers: datasets

Let's take an example by running a convnets on of image of dimension $32 \times 32 \times 3$.

- Input Layers:** It's the layer in which we give input to our model. In CNN, Generally, the input will be an image or a sequence of images. This layer holds the raw input of the image with width 32, height 32, and depth 3.

- Convolutional Layers:** This is the layer, which is used to extract the feature from the input dataset. It applies a set of learnable filters known as the kernels to the input images. The filters/kernels are smaller matrices usually 2×2 , 3×3 , or 5×5 shape. it slides over the input image data and computes the dot product between kernel weight and the corresponding input image patch. The output of this layer is referred as feature maps. Suppose we use a total of 12 filters for this layer we'll get an output volume of dimension $32 \times 32 \times 12$.

- Activation Layer:** By adding an activation function to the output of the preceding layer, activation layers add nonlinearity to the network. it will apply an element-wise activation function to the output of the convolution layer. Some common activation functions are RELU: $\max(0, x)$, Tanh, Leaky RELU, etc. The volume remains unchanged hence output volume will have dimensions $32 \times 32 \times 12$.




- put of each class into the probability score of each class. **Pooling layer:** This layer is periodically inserted in the convnets and its main function is to reduce the size of volume which makes the computation fast reduces memory and also prevents overfitting. Two common types of pooling layers are max pooling and average pooling. If we use a max pool with 2 x 2 filters and stride 2, the resultant volume will be of dimension 16x16x12.

- Flattening:** The resulting feature maps are flattened into a one-dimensional vector after the convolution and pooling layers so they can be passed into a completely linked layer for categorization or regression.

- Fully Connected Layers:** It takes the input from the previous layer and computes the final classification or regression task.

- Output Layer:** The output from the fully connected layers is then fed into a logistic function for classification tasks like sigmoid or softmax which converts the out




Recurrent Neural Network

we will introduce a new variation of neural network which is the Recurrent Neural Network also known as (RNN) that works better than a simple neural network when data is sequential like Time-Series data and text data.

What is Recurrent Neural Network (RNN)?


Recurrent Neural Network(RNN) is a type of Neural Network where the output from the previous step is fed as input to the current step. In traditional neural networks, all the inputs and outputs are independent of each other. Still, in cases when it is required to predict the next word of a sentence, the previous words are required and hence there is a need to remember the previous words. Thus RNN came into existence, which solved this issue with the help of a Hidden Layer. The main and most important feature of RNN is its Hidden state, which remembers some information about a sequence. The state is also referred to as Memory State since it remembers the previous input to the network. It uses the same parameters for each input as it performs the same task on all the inputs or hidden layers to produce the output. This reduces the complexity of parameters, unlike other neural networks.



How RNN differs from Feedforward Neural Network?

Artificial neural networks that do not have looping nodes are called feed forward neural networks. Because all information is only passed forward, this kind of neural network is also referred to as a multi-layer neural network.


Information moves from the input layer to the output layer – if any hidden layers are present – unidirectionally in a feedforward neural network. These networks are appropriate for image classification tasks, for example, where input and output are independent. Nevertheless, their inability to retain previous inputs automatically renders them less useful for sequential data analysis.



Recurrent Neuron and RNN Unfolding

The fundamental processing unit in a Recurrent Neural Network (RNN) is a Recurrent Unit, which is not explicitly called a “Recurrent Neuron.” This unit has the unique ability to maintain a hidden state, allowing the network to capture sequential dependencies by remembering previous inputs while processing. [Long Short-Term Memory \(LSTM\)](#) and [Gated Recurrent Unit \(GRU\)](#) versions improve the RNN’s ability to handle long-term dependencies.

Recurrent Neuron




Types Of RNN

There are four types of RNNs based on the number of inputs and outputs in the network.

2. One to One
3. One to Many
4. Many to One
5. Many to Many


One to One

This type of RNN behaves the same as any simple Neural network it is also known as Vanilla Neural Network. In this Neural network, there is only one input and one output.




One To Many

In this type of RNN, there is one input and many outputs associated with it. One of the most used examples of this network is Image captioning where given an image we predict a sentence having Multiple words.




Many to One

In this type of network, Many inputs are fed to the network at several states of the network generating only one output. This type of network is used in the problems like sentimental analysis. Where we give multiple words as input and predict only the sentiment of the sentence as output.



Many to Many

In this type of neural network, there are multiple inputs and multiple outputs corresponding to a problem. One Example of this Problem will be language translation. In language translation, we provide multiple words from one language as input and predict multiple words from the second language as output.



Recurrent Neural Network Architecture

RNNs have the same input and output architecture as any other deep neural architecture. However, differences arise in the way information flows from input to output. Unlike Deep neural networks where we have different weight matrices for each Dense network in RNN, the weight across the network remains the same. It calculates state hidden state H_i for every input X_i . By using the following formulas:

$$h = \sigma(UX + Wh_{-1} + B)$$


$$Y = O(Vh + C)$$

Hence

$$Y = f(X, h, W, U, V, B, C)$$

Here S is the State matrix which has element s_i as the state of the network at timestep i

The parameters in the network are W, U, V, c, b which are shared across timestep



How does RNN work?

The Recurrent Neural Network consists of multiple fixed activation function units, one for each time step. Each unit has an internal state which is called the hidden state of the unit. This hidden state signifies the past knowledge that the network currently holds at a given time step. This hidden state is updated at every time step to signify the change in the knowledge of the network about the past. The hidden state is updated using the following recurrence relation:-

The formula for calculating the current state:

$$h_t = f(h_{t-1}, x_t)$$

where,

- h_t -> current state
- h_{t-1} -> previous state
- x_t -> input state

Formula for applying Activation function(tanh)

$$h_t = \tanh(W_{hh}h_{t-1} + W_{hx}x_t)$$

where,

- W_{hh} -> weight at recurrent neuron
- W_{hx} -> weight at input neuron

The formula for calculating output:

$$y_t = W_{hy}h_t$$

- Y_t -> output
- W_{hy} -> weight at output layer

These parameters are updated using Backpropagation. However, since RNN works on sequential data here we use an updated backpropagation which is known as Backpropagation through time.

Optimizers

Deep learning, a subset of machine learning, tackles intricate tasks like speech recognition and text classification. Comprising components such as activation functions, input, output, hidden layers, and loss functions, deep learning models aim to generalize data and make predictions on unseen data. To optimize these models, various algorithms, known as optimizers, are employed. Optimizers adjust model parameters iteratively during training to minimize a loss function, enabling neural networks to learn from data. This guide delves into different optimizers used in deep learning, discussing their advantages, drawbacks, and factors influencing the selection of one optimizer over another for specific applications. Common optimizers include Stochastic Gradient Descent (SGD), Adam, and RMSprop, each employing specific update rules, learning rates, and momentum for refining model parameters. Optimizers play a pivotal role in enhancing accuracy and speeding up the training process, shaping the overall performance of deep learning models.

Need for Optimizers in Deep Learning

Choosing an appropriate optimizer for a deep learning model is important as it can greatly impact its performance. Optimization algorithms have different strengths and weaknesses and are better suited for certain problems and architectures.

For example, stochastic gradient descent is a simple and efficient optimizer that is widely used, but it may need help to converge on problems with complex, non-convex loss functions. On the other hand, Adam is a more sophisticated optimizer that combines the ideas of momentum and adaptive learning rates and is often considered one of the most effective optimizers in deep learning.

Here are a few pointers to keep in mind when choosing an optimizer:

- Understand the problem and model architecture, as this will help you determine which optimizer is most suitable
- Experiment with different optimizers in deep learning to see which one works best for your problem
- Adjust the hyperparameters of the optimizer, such as the learning rate, to see if it improves performance
- Remember that the optimizer's choice is not the only factor affecting model performance.
- Other important factors include the choice of architecture, the quality of the data, and the amount of data available.

Types of Optimizers

Many types of optimizers are available for training machine learning models, each with its own strengths and weaknesses. Some optimizers are better suited for certain types of models or data, while others are more general-purpose.


This section will briefly overview the most commonly used optimizers, starting with the simpler ones and progressing to the more complex ones.

Gradient Descent

Gradient descent is a simple optimization algorithm that updates the model's parameters to minimize the loss function. We can write the basic form of the algorithm as follows:

$$\theta = \theta - \alpha \cdot \nabla_{\theta} L(\theta)$$

where θ is the model parameter, $L(\theta)$ is the loss function, and α is the learning rate.



Pros:

- Simple to implement.
- Can work well with a well-tuned learning rate.

Cons:


- It can converge slowly, especially for complex models or large datasets.
- Sensitive to the choice of learning rate.

Stochastic Gradient Descent

Stochastic gradient descent (SGD) is a variant of gradient descent that involves updating the parameters based on a small, randomly-selected subset of the data (i.e., a "mini-batch") rather than the full dataset. We can write the basic form of the algorithm as follows:

$$\theta = \theta - \alpha \cdot \nabla \theta L(\theta; x(i); y(i))$$

where $(x(i), y(i))$ is a mini-batch of data.



Pros:

- It can be faster than standard gradient descent, especially for large datasets.
- Can escape local minima more easily.

Cons:

- It can be noisy, leading to less stability.
- It may require more hyperparameter tuning to get good performance.


Stochastic Gradient Descent with Momentum

SGD with momentum is a variant of SGD that adds a "momentum" term to the update rule, which helps the algorithm move in the same direction even if the local gradient is noisy. The momentum term is typically set to a value between 0 and 1.

$$v = \beta \cdot v + (1 - \beta) \cdot \nabla \theta_i$$

$$\theta = \theta - \alpha \cdot v$$

where v is the momentum vector and β is the momentum hyperparameter



Pros:


- It can help the optimizer to move more efficiently through "flat" regions of the loss function.
- It can help to reduce oscillations and improve convergence.

Cons:

- Can overshoot good solutions and settle for suboptimal ones if the momentum is too high.
- Requires tuning of the momentum hyperparameter.

Mini-Batch Gradient Descent

Mini-batch gradient descent is similar to SGD, but instead of using a single sample to compute the gradient, it uses a small, fixed-size "mini-batch" of samples. The update rule is the same as for SGD, except that the gradient is averaged over the mini-batch. This can reduce noise in the updates and improve convergence.

**Pros:**


- It can be faster than standard gradient descent, especially for large datasets.
- Can escape local minima more easily.
- Can reduce noise in updates, leading to more stable convergence.

Cons:

- Can be sensitive to the choice of mini-batch size.

How Do Optimizers Work in Deep Learning?

Optimizers in deep learning adjust the model's parameters to minimize the loss function. The loss function measures how well the model can make predictions on a given dataset, and the goal of training a model is to find the set of model parameters that yields the lowest possible loss.



The optimizer uses an optimization algorithm to search for the parameters that minimize the loss function. The optimization algorithm uses the gradients of the loss function to the model parameters to determine the direction in which we should adjust the parameters.

The gradients are computed using backpropagation, which involves applying the chain rule to compute the gradients of the loss function to each of the model parameters.

The optimization algorithm then adjusts the model parameters to minimize the loss function. This process is repeated until the loss function reaches a minimum or the optimizer reaches the maximum number of allowed iterations.

Conclusion

Optimizers in deep learning are essential, as they adjust the model's parameters to minimize the loss function.

In general, the choice of which optimization algorithm to use will depend on the specific characteristics of the problem, such as the dataset's size and the model's complexity.

It is important to consider each algorithm's pros and cons carefully and tune any relevant hyperparameters to achieve the best possible performance.

Overall, understanding the role of optimizers in deep learning and the various available algorithms is essential for anyone looking to build and train effective machine learning models.

ASSIGNMENT – UNIT I

1. Design a CNN architecture for detecting objects in images.(K6)
2. Evaluate the trade-offs between convergence speed and stability for different optimizers(K5)
3. Analyze the impact of learning rate on the effectiveness of backpropagation.(K4)
4. Design a deep neural network architecture for image recognition(K6)
5. How would you apply an RNN to a time series prediction task?(K3)



PART A- UNIT-1

1. Define what a generative model is in the context of machine learning.

Generative" describes a class of statistical models that contrasts with discriminative models. Informally: Generative models can generate new data instances. Discriminative models discriminate between different kinds of data instances.

2. What are the challenges associated with training deep neural networks compared to shallow networks?

One challenge is that they can be more difficult to train than shallow networks, due to the increased number of parameters and the risk of overfitting. Another challenge is that deep networks can be computationally expensive to train and require a lot of data to achieve good performance.

3. What do you mean by DNN?

Deep neural networks (DNN) can be defined as ANNs with additional depth, that is, an increased number of hidden layers between the input and the output layers.

4. What is a perceptron?

A machine-based algorithm used for supervised learning of various binary sorting tasks is called Perceptron. Furthermore, Perceptron also has an essential role as an Artificial Neuron or Neural link in detecting certain input data computations in business intelligence

5. Define backpropagation in the context of neural networks.

Backpropagation is an algorithm used in artificial intelligence and machine learning to train artificial neural networks through error correction. The computer learns by calculating the loss function, or the difference between the input you provided and the output it produced.

6. How many convolutional layers are in CNN?

A CNN typically has three layers: a convolutional layer, a pooling layer, and a fully connected layer.

7. What are convolutional neural networks?

A convolutional neural network (CNN or ConvNet) is a network architecture for deep learning that learns directly from data. CNNs are particularly useful for finding patterns in images to recognize objects, classes, and categories. They can also be quite effective for classifying audio, time-series, and signal data.

8. How would you apply an RNN to a time series prediction task?

A typical structure for time series prediction comprises an input layer, one or more hidden layers with LSTM or GRU cells and an output layer. Selecting the optimal number of layers and neurons: This is a critical step in building the RNN model.

9. Define recurrent neural networks (RNNs)


A recurrent neural network (RNN) is a deep learning model that is trained to process and convert a sequential data input into a specific sequential data output.

10. What is an optimizer in the context of neural networks?

Optimizers are algorithms or methods used to change the attributes of the neural network such as weights and learning rate to reduce the losses. Optimizers are used to solve optimization problems by minimizing the function

PART B- UNIT 1

1. Apply a CNN to a standard image classification task.K3
- 2.Explain how RNNs process sequential data.K2
- 3.Compare the performance of RNNs on different types of sequential data.K4
- 4.Explain the difference between gradient descent and stochastic gradient descent.K2
- 5.Analyze the convergence behavior of different optimizers.K4
- 6.Evaluate the trade-offs between speed and accuracy for various optimizers.K5
- 7.Design a multi-layer perceptron to solve a complex classification task.K6
- 8.Analyze the impact of network depth on the training processK4
- 9.Apply image transformation techniques to preprocess a dataset.K3
- 10.Demonstrate how to evaluate the performance of a generative model.K3



R.M.K.
GROUP OF
INSTITUTIONS

SUPPORTIVE ONLINE COURSES – UNIT I

✿ https://onlinecourses.swayam2.ac.in/imb24_mg116/preview

✿ [Introduction to Generative AI Course by Google Cloud | Coursera](#)


✿ <https://www.coursera.org/learn/generative-ai-for-everyone>

✿ [Generative AI Fundamentals | Coursera](#)




Real Time Applications in Day to Day life and to Industry

1.Face detection: CNNs have been used to detect faces within images. The network takes an image as the input and produces a set of values that represent characteristics of faces or facial features at different parts of the image. CNN has shown improved accuracy over previous algorithms, identifying faces 97% of the time according to various research papers. CNN can also be used to help reduce the amount of distortion in a face. CNNs can detect features such as eyes, nose, and mouth with great accuracy reducing distortions from things like angles or shadows on faces.



2.Object detection: CNN has been applied to object recognition across images by classifying objects based on shapes and patterns found within an image. CNN models have been created that can detect a wide range of objects from everyday items such as food, celebrities, or animals to more unusual ones including dollar bills and guns. Object detection is performed using techniques such as semantic or instance segmentation. CNNs have been used to localize and identify objects within images as well as create different views of those objects such as for use in drones or self-driving cars.




3.Speech Recognition

RNNs can be used for predicting phonetic segments considering sound waves from a medium as an input source .The set of inputs consists of phoneme or acoustic signals from an audio which are processed in a proper manner and taken as inputs. The RNN network will compute the phonemes and then produce a phonetic segment along with the likelihood of output.The steps used in speech recognition are as follows:-

The input data is first processed and recognized through a neural network. The result consists of a varied collection of input sound waves.


The information contained in the sound wave is further classified by intent and through keywords related to the query.

Then input sound waves are classified into phonetic segments and are pieced together into cohesive words using a RNN application. The output consists of a pattern of phonetic segments put together into a singular whole in a logical manner.



4. Language Modelling and Generating Text

Taking a sequence of words as input, we try to predict the possibility of the next word. This can be considered to be one of the most useful approaches for translation since the most likely sentence would be the one that is correct. In this method, the probability of the output of a particular time-step is used to sample the words in the next iteration.



PRESCRIBED TEXT BOOKS AND REFERENCE BOOKS

TEXT BOOKS:

1. Ben Auffarth, Generative AI with LangChain, Packt Publishing, 2023.
2. Amit Bahree, Generative AI in Action, Manning Publication, First Edition, 2023.

REFERENCES:

1. David Foster, Generative Deep Learning, 2nd Edition, O'Reilly Media, 2023.
2. Numa Dhamani and Maggie Engler, Introduction to Generative AI, Manning Publication, First Edition, 2024.
3. Valentina Alto, Modern Generative AI with ChatGPT and OpenAI Models, Packt publications, 2024.



R.M.K.
GROUP OF
INSTITUTIONS

MINI PROJECT SUGGESTIONS

- ❁ **TEAM 1:** Neural Network Project on Automatic Music Generation System
- ❁ **TEAM 2:** Neural Network Project on Handwriting Recognition Tool using Autoencoders
- ❁ **TEAM 3:** Neural Network Project on Stock Market Value Prediction System using RNN
- 1. **TEAM 4:** Neural Network Project to Build a Vehicle Security System
- ❁ **TEAM 5:** Neural Network Project on Gender Recognition Systems



TOUGH QUESTIONS SIMILAR TO GATE

1. Which statement is false about GANs vs VAEs?

- A) GANs minimize Jensen-Shannon divergence, VAEs minimize KL divergence.
- B) VAEs provide explicit likelihood, GANs do not.
- C) GANs are prone to mode collapse, VAEs generate blurry samples.
- D) VAEs use adversarial training, GANs use reparameterization.

Ans: D

2. A 3×3 image undergoes 90° rotation followed by horizontal flipping.

The final pixel at (0,0) was originally at:

- A) (0, 2)
- B) (2, 0)
- C) (2, 2)
- D) (0, 0)

Ans: B.

3. Which is not a solution for mode collapse?

- A) Mini-batch discrimination
- B) Adding noise to discriminator inputs
- C) Using Wasserstein loss
- D) Increasing latent space dimension

Ans: D



TOUGH QUESTIONS SIMILAR TO GATE

4. What is the output of this perceptron?

```
import numpy as np
```

```
def perceptron(x, w, b):
```

```
    return 1 if np.dot(x, w) + b > 0 else 0
```

```
x = [1, -2]; w = [3, 4]; b = -5
```

A) 0

B) 1

C) Runtime error

D) -1

Ans: A

5. For a network with loss

$$L = (y - \hat{y})^2, \hat{y} = \sigma(wx + b), \text{ what is } \frac{\partial L}{\partial w}?$$

A) $2(y - \hat{y}) \cdot \sigma'(wx + b) \cdot x$

B) $-2(y - \hat{y}) \cdot \sigma'(wx + b) \cdot x$

C) $2(y - \hat{y}) \cdot \sigma'(wx + b) \cdot (1 - \sigma(wx + b)) \cdot x$

D) $-2(y - \hat{y}) \cdot x$

Ans: C (Chain rule: $\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z} \cdot \frac{\partial z}{\partial w}$, where $z = wx + b$.)

TOUGH QUESTIONS SIMILAR TO GATE

6. Gives the tensorflow code:

```
conv = tf.keras.layers.Conv2D(filters=32, kernel_size=3, strides=2,  
padding='same')
```

What is the output shape for input (None, 64, 64, 3)?

- A) (None, 32, 32, 32)
- B) (None, 32, 32, 3)
- C) (None, 64, 64, 32)
- D) (None, 64, 64, 3)

Ans: A.

7. In an RNN, the hidden state

$$h_t = \tanh(W_{xh}x_t + W_{hh}h_{t-1} + b).$$

If h_{t-1} has 128 units and x_t has 64. What is

W_{hh} 's shape?

- A) (64, 128)
- B) (128, 128)
- C) (64, 64)
- D) (128, 64)

Ans : B8.

TOUGH QUESTIONS SIMILAR TO GATE

8. Which is true about Adam?

- A)** Combines momentum and RMSProp with bias correction.
- B) Uses constant learning rate for all parameters
- C) Does not track second moments of gradients.
- D) Is equivalent to SGD with Nesterov acceleration.

Ans: A

9. Gradients vanish in a 10-layer ReLU network. Which change helps?

- A) Replace ReLU with LeakyReLU.
- B) Reduce learning rate.
- C) Use SGD instead of Adam.
- D) Remove batch normalization.

Ans: A

10. A 5×5 image convolved with a 3×3 kernel (stride=1, no padding) produces output of size:

- A) 3×3
- B) 5×5
- C) 4×4
- D) 2×2

Ans : A

THANK YOU

Disclaimer:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

R.M.K GROUP OF ENGINEERING INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS

R.M.K GROUP OF INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS



Please read this disclaimer before proceeding:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

Digital Course Material
22AI005
Introduction to Generative AI



Department : CSE(CYBER SECURITY)

Batch/Year : 2022-2026/IV

Created by : Ms G.K.MONICA

AP/CSE(CS)

Date : 07-07-2025

TABLE OF CONTENTS

S.NO	CONTENTS	SLIDE NO.
1	COURSE OBJECTIVES	6
2	PRE REQUISITES (COURSE NAMES WITH CODE)	7
3	SYLLABUS (WITH SUBJECT CODE, NAME, LTPC DETAILS)	8
4	COURSE OUTCOMES (6)	9
5	CO- PO/PSO MAPPING	10
6	LECTURE PLAN –UNIT 2	12
7	ACTIVITY BASED LEARNING –UNIT 2	14
8	TEST YOURSELF	15
9	LECTURE NOTES	16
10	ASSIGNMENT	29
11	PART A Q s (WITH K LEVEL AND CO)	30
12	PART B Q s (WITH K LEVEL AND CO)	31
13	SUPPORTIVE ONLINE CERTIFICATION COURSES	33
14	REAL TIME APPLICATIONS IN DAY - TO - DAY LIFE AND TO INDUSTRY	34
15	PRESCRIBED TEXT BOOKS & REFERENCE BOOKS	37
16	MINI PROJECT SUGGESTIONS	38
17	TOUGH QUESTIONS SIMILAR TO GATE	39

COURSE OBJECTIVES


The Course will enable learners to:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.



PRE REQUISITES

✿ PRE-REQUISITE CHART



OBJECTIVES:

- ✿ To understand the basic concepts of Generative AI.
- ✿ To build Generative AI systems to generate images.
- ✿ To understand the concept used in Generative AI Models.
- ✿ To use various Generative AI models.
- ✿ To compare and use the various Large Language Models.
- ✿ To understand the basics of Prompt Engineering.

UNIT I INTRODUCTION**9**

Generative Models – Image transformation – Challenges - Deep Neural Networks – Perceptron – back propagation – CNN – RNN – Optimizer.

UNIT II IMAGE GENERATION**9**

Creating encodings of images – variational objective – Inverse Autoregressive flow – Importing CIFAR – Creating the network from TensorFlow 2.

UNIT III GENERATIVE ADVERSARIAL NETWORKS**9**

Generative Adversarial Networks – Vanilla GAN – Improved GANs – Progressive GAN – Challenges – Paired style transfer – Unpaired style transfer – Deepfakes – Modes of operation – key feature set – High level flow – Replacement – Re-enactment.

UNIT IV LARGE LANGUAGE MODELS**9**

Overview of LLMs - Transformers – GPT – Types of LLMs – Key concepts – other Transformers – T5 – Generative Pre-Training Models – Multi-modal Models – DALL.E 2

UNIT V PROMPT ENGINEERING**9**

Basics – In-Context Learning – In-Context Prompting – Techniques – Image Prompting – Prompt Hijacking – Challenges.

TOTAL: 45 PERIODS

COURSE OUTCOME

Course Code	Course Outcome Statement	Cognitive/ Affective Level of the Course Outcome	Expected Level of Attainment
Course Outcome Statements in Cognitive Domain			
C305.1	Elaborate the basic concepts of Generative AI	Understand K2	60%
C305.2	Build Generative AI systems to generate images	Analyse K4	60%
C305.3	Apply the concepts used in Generative AI Models	Apply K3	60%
C305.4	Use various Generative AI models.	Apply K3	60%
C305.5	Compare and use the various Large Language Models	Analyse K4	60%
C305.6	Analyze the basics of Prompt Engineering.	Apply K3	60%
Course Outcome Statements in Affective domain			
C305.7	Attend the classes regularly	Respond (A2)	95%
C305.8	Submit the Assignments regularly.	Respond (A2)	95%
C305.9	Participation in Seminar/Quiz/ Group Discussion/ Collaborative learning and content beyond syllabus	Valuing (A3)	95%

CO-PO/PSO MAPPING

Correlation Matrix of the Course Outcomes to Programme Outcomes and Programme Specific Outcomes Including Course Enrichment Activities

Course Outcomes (Cos)		Programme Outcomes (POs), Programme Specific Outcomes (PSOs)														
		PO 1	PO 2	PO 3	PO 4	PO 5	PO 6	PO 7	PO 8	PO 9	PO 10	PO 11	PO 12	PSO 1	PSO 2	PSO 3
		K 3	K4	K 5	K5	K3 /K 5	A2	A3	A 3	A3	A3	A3	A2	K3	K3	K3
C305.1	K 3	3	2	1	1	3						3		3	3	3
C305.2	K 4	3	3	2	2	3						3		3	3	3
C305.3	K 2	2	1									2		3	3	3
C305.4	K 3	3	2	1	1	3						3		3	3	3
C305.5	K 4	3	3	2	2	3						3		3	3	3
C305.6	K 3	3	2	1	1	3						3		2	2	2
C305.7	A 2												3			
C305.8	A 2								2	2	2		3			
C305.9	A 3						3	3		3	3		3			
C305		3	3	2	2	3	1	1	1	3	3	3	3	3	3	3



UNIT II

IMAGE GENERATION

LECTURE PLAN – UNIT II

UNIT II IMAGE GENERATION							
Sl. No	TOPIC	NO OF PERIODS	PROPOSED LECTURE	ACTUAL LECTURE	PERTAINING CO(s)	TAXONOMY LEVEL	MODE OF DELIVERY
			PERIOD	PERIOD			
1	Creating encodings of images	1			CO2	K2	PPT
2	variational objective	1			CO2	K2	PPT
3	Inverse Autoregressive flow	1			CO2	K2	PPT
4	Importing CIFAR	1			CO2	K3	PPT
5	Creating the network from TensorFlow 2	1			CO2	K3	PPT

LECTURE PLAN – UNIT II

✿ ASSESSMENT COMPONENTS

- ✿ AC 1. Unit Test
- ✿ AC 2. Assignment
- ✿ AC 3. Course Seminar
- ✿ AC 4. Course Quiz
- ✿ AC 5. Case Study
- ✿ AC 6. Record Work
- ✿ AC 7. Lab / Mini Project
- ✿ AC 8. Lab Model Exam
- ✿ AC 9. Project Review

MODE OF DELEIVERY

- MD 1. Oral presentation
- MD 2. Tutorial
- MD 3. Seminar
- MD 4 Hands On
- MD 5. Videos
- MD 6. Field Visit



R.M.K.
GROUP OF
INSTITUTIONS

ACTIVITY BASED LEARNING – UNIT II

QUIZ- LINK

Unit II:

Inverse Autoregressive flow

<https://testbook.com/objective-questions/mcq-on-type-of-flow--5eea6a0c39140f30f369e18e>

Encodings of images

<https://quizgecko.com/learn/image-compression-and-encoding-kk19ep>

CIFAR

<https://www.cs.toronto.edu/~kriz/cifar.html>

TensorFlow 2:

<https://www.acmecollinsschool.com/tensorflow-2-mcqs-questions.html>

TEST YOURSELF

1.Which of the following techniques is used to ensure that the latent space in a Variational Autoencoder is well-behaved and follows a normal distribution?

- A. Reconstruction loss
- B. Cross-entropy loss
- C. KL Divergence loss
- D. Dropout regularization

2..What is the primary purpose of creating encodings of images in deep learning?

- A. To improve image resolution
- B. To reduce the dimensionality of the data while preserving important information
- C. To change the color scheme of the image
- D. To segment the image into different regions

3.What is the primary purpose of using Inverse Autoregressive Flow (IAF) in variational inference?

- A. To reduce the computational complexity of training deep neural networks
- B. To enhance the expressiveness of the approximate posterior distribution
- C. To perform dimensionality reduction on input data
- D. To visualize high-dimensional data

4.TensorFlow Data API does not support text files?

- A. True
- B. False

5.Using TensorFlow Data API we can do the following:

- A. Batching
- B. Prefetching
- C. Gradient Descent
- D. All of the above

UNIT II

Creating encodings of images

Creating encodings of images involves transforming images into a compact, fixed-size representation (or vector) that captures the essential features and information. These encodings are used in various applications such as image retrieval, classification, clustering, and as inputs to other machine learning models.

Overview

- 1. Feature Extraction:** Extract meaningful features from the image using pre-trained convolutional neural networks (CNNs).
- 2. Dimensionality Reduction (Optional):** Reduce the dimensionality of the extracted features to create a more compact encoding.
- 3. Storage/Usage:** Use these encodings for downstream tasks or store them for later use.

Steps to Create Image Encodings

1. Feature Extraction

The most common approach is to use a pre-trained CNN, such as VGG16, ResNet, or Inception, to extract features from images. These models, trained on large datasets like ImageNet, can effectively capture general features.

Using Pre-trained Models:

```
import torch

import torch.nn as nn

from torchvision import models, transforms

from PIL import Image

# Load the pre-trained ResNet model

model = models.resnet50(pretrained=True)

# Remove the last fully connected layer

model = nn.Sequential(*list(model.children())[:-1])

# Set the model to evaluation mode

model.eval()
```

Image transformations

```
preprocess = transforms.Compose
```

```
([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])
```

```
def get_image_encoding(image_path):
```

Load the image

```
image = Image.open(image_path)
image = preprocess(image)
image = image.unsqueeze(0) # Add batch dimension
```

Extract features

```
with torch.no_grad():
    encoding = model(image)
```

Flatten the encoding

```
encoding = encoding.view(encoding.size(0), -1)
return encoding.numpy()
```

Example usage

```
image_path = 'path_to_image.jpg'
encoding = get_image_encoding(image_path)
print(encoding)
```

Dimensionality Reduction (Optional)

Depending on the application, the dimensionality of the extracted features might be reduced. Techniques like PCA (Principal Component Analysis) or t-SNE (t-Distributed Stochastic Neighbor Embedding) are commonly used.

Using PCA for Dimensionality Reduction:

```
From sklearn.decomposition import PCA
```

Suppose encodings is a list of image encodings obtained from the above step

```
encodings = [get_image_encoding(img_path) for img_path in list_of_image_paths]
```

```
encodings = np.vstack(encodings) # Stack to a single array
```

Apply PCA to reduce dimensionality

```
pca = PCA(n_components=128) # Example: Reduce to 128 dimensions
```

```
reduced_encodings = pca.fit_transform(encodings)
```

Storage/Usage

The encodings can be stored in a file or database for later retrieval or used directly in downstream machine learning tasks.

Storing Encodings:

```
import numpy as np
```

Save to a file

```
np.save('image_encodings.npy', reduced_encodings)
```

Load from a file

```
loaded_encodings = np.load('image_encodings.npy')
```

Applications

- 1. Image Retrieval:** Encodings can be used to find similar images by comparing their encodings using distance metrics like Euclidean distance.
- 2. Classification:** Encodings can serve as inputs to classifiers (e.g., SVM, random forest) for tasks such as image classification.
- 3. Clustering:** Encodings can be clustered using algorithms like K-means to group similar images.
- 4. Anomaly Detection:** Encodings can help in identifying outliers or anomalies in image datasets.

variational objective

The objective of a Variational Autoencoder (VAE) is to learn a probabilistic latent space representation of the input data such that it can generate new data points similar to the input data. This is achieved through two main components: the encoder (which maps the input data to a latent space) and the decoder (which maps the latent space back to the data space).

The training objective of a VAE combines two loss terms:

- 1. Reconstruction Loss:** This term measures how well the decoder reconstructs the input data from the latent representation. It is typically measured using Mean Squared Error (MSE) or Binary Cross-Entropy (BCE) between the original input data and the reconstructed data. The reconstruction loss ensures that the latent space captures the important features of the input data.
- 2. KL Divergence Loss:** This term measures how close the learned latent distribution $q(z|x)$ (produced by the encoder) is to the prior distribution $p(z)$. The prior is usually a standard normal distribution $N(0, I)$. The KL Divergence encourages the latent representations to be similar to this prior distribution, which helps in regularizing the latent space and ensuring that the generated data points are diverse and meaningful.

Mathe

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - \text{KL}(q_{\phi}(z|x) || p(z))$$

Where

- Θ are the parameters of the decoder.
- Φ are the parameters of the encoder.
- x is the input data.
- z is the latent representation.
- $p_{\theta}(x|z)$ is the likelihood of the data given the latent variable.
- $q_{\phi}(z|x)$ is the approximate posterior distribution (encoder).
- $p(z)$ is the prior distribution over the latent variables.

Explanation

$$(\mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)]):$$

- **Reconstruction** This term ensures that the decoded output (reconstructed input) is as close as possible to the original input. It forces the decoder to learn to generate realistic data from the latent representations $(\text{KL}(q_{\phi}(z|x) || p(z)))$:
- **KL Divergence Term** This term regularizes the latent space by ensuring that the distribution of the latent variables z stays close to the prior distribution $p(z)$, usually a standard Gaussian. This regularization helps in generating new, diverse data samples and prevents overfitting.

By balancing these two terms, the VAE ensures that the latent space is both informative for reconstructing the data and well-regularized to allow for meaningful sampling and generation of new data points.

Inverse Autoregressive flow

Inverse Autoregressive Flow (IAF) is a powerful technique used in the field of probabilistic machine learning, particularly for variational inference and generative modeling. It allows for the construction of complex, flexible distributions by transforming simple base distributions through a sequence of invertible transformations.

Autoregressive Models: Autoregressive models are a class of models where each output depends on the previous outputs. For a sequence $\mathbf{x} = (x_1, x_2, \dots, x_D)$, an autoregressive model specifies the joint distribution as a product of conditional distributions:

$$p(\mathbf{x}) = \prod_{i=1}^D p(x_i | x_{1:i-1})$$

Normalizing Flows: Normalizing flows transform a simple base distribution $p(\mathbf{z})$ (e.g., Gaussian) into a more complex distribution $p(\mathbf{x})$ using a sequence of invertible and differentiable transformations:

Inverse Autoregressive Flow

Definition: IAF is a specific type of normalizing flow where the transformations are designed such that the inverse transformations are autoregressive. This is useful because autoregressive transformations allow efficient sampling and density evaluation.

Transformation Function: IAF uses a transformation function f parameterized as follows:

$$\mathbf{x} = \mu(\mathbf{z}) + \sigma(\mathbf{z}) \odot \mathbf{z}$$

where:

- $\mu(\mathbf{z})$ is the mean function.
- $\sigma(\mathbf{z})$ is the scale function.
- $\odot$ denotes element-wise multiplication.

Sampling Process:

1. Start with a latent variable $\mathbf{z}$ from a simple base distribution.
2. Apply the transformation to obtain the sample $\mathbf{x}$.

Density Calculation:

Using the change of variables formula, the density $p(\mathbf{x})$ can be computed as:

$$p(\mathbf{x}) = p(\mathbf{z}) \left| \det \left(\frac{\partial \mathbf{z}}{\partial \mathbf{x}} \right) \right|$$

Since $\mathbf{z}$ is a simple base distribution (e.g., Gaussian), $p(\mathbf{z})$ is easy to compute. The determinant of the Jacobian $\left| \det \left(\frac{\partial \mathbf{z}}{\partial \mathbf{x}} \right) \right|$ can be efficiently computed using the autoregressive property.

Key Properties

1.Invertibility: The transformation is designed to be invertible, ensuring that both sampling and density evaluation are feasible.

2.Expressiveness: IAF can represent complex distributions due to the flexibility of the transformation functions μ and σ .

3.Autoregressive Structure: The autoregressive nature allows efficient computation of the Jacobian determinant, crucial for density estimation.

Applications

1.Variational Inference: IAF is often used in Variational Autoencoders (VAEs) to approximate the posterior distribution. The flexibility of IAF allows for better approximation of complex posterior distributions.

2.Generative Modeling: In models like PixelCNN, IAF helps in generating high-quality samples by modeling complex distributions of pixel values.

Implementation Considerations

•**Neural Networks for μ and σ :** Typically, neural networks parameterize the functions μ and σ .

•**Efficiency:** While IAF can be computationally intensive due to the need for sequential transformations, it is still efficient compared to other complex distribution modeling techniques.

Creating the network from TensorFlow 2

A neural network comprises several nodes in each layer of the neural network architecture which are interconnected. They are important for learning new features from instances of the data. Thus, they prove to be more accurate than most machine learning algorithms like logistic regression, support vector machine, etc. They can be used for processing images and identifying the objects within them. There are many other use cases of neural networks where they have proved to be very accurate. In this article, we'll be focusing on TensorFlow neural networks and how to build them.

What is TensorFlow?

TensorFlow is an open-source library developed in Python for the implementation of machine learning and artificial intelligence. It was initially developed for the Python language but the support has now extended to other languages like Java, C++, and JavaScript. The library was built by Google Brain with focus on implementing deep neural networks.

Some of the important features of TensorFlow are:

1. Automatic differentiation: A process that calculates the gradient vector of each of the parameters in the model that help in optimizing the performance of the model, such as backpropagation.
2. Eager execution: A mode that directly executes the code rather than adding to a computational graph that executes the code at a later stage.
3. Distribute: An API that distributes the computation of models on various devices which results in speeding up the computation process.
4. Losses: Functions that are required for assessing a trained model. They compare the model's output with the expected output and compute the error loss. A few popular loss functions are binary cross entropy and mean squared error.

Setting up TensorFlow

To install the latest version of TensorFlow, the following system requirements must be met:

1. 64-bit operating system
2. Python version 3.7 or later
3. Windows 7 or later
4. macOS 10.12.6 (Sierra) or later (no GPU support)
5. Ubuntu 16.04 or later

The TensorFlow library can be installed using the following commands:

```
# Upgrade pip to latest version
```

```
pip install --upgrade pip
```

```
#Installation using pip command
```

```
pip install tensorflow
```

There is another way of using TensorFlow too: with Docker. The only requirement is that the docker should be installed on the PC. You can run the following command in the terminal:

```
docker pull tensorflow/tensorflow:latest # Download latest stable image
```

```
docker run -it -p 8888:8888 tensorflow/tensorflow:latest-jupyter # Start Jupyter server
```

Developing a TensorFlow neural network

This section of the article will discuss some basic steps for developing a neural network using TensorFlow. Before this, we will look at basic terminologies related to convolutional neural networks (CNNs):

1. CNNs are a type of artificial neural network (ANN) which are very popular among machine learning practitioners. They consist of several layers, such as the convolutional layer, pooling layer, and output layer.
2. The convolutional layer is added to the architecture to extract useful features of image input.
3. The feature extraction in the previous step is achieved with the help of kernels. Kernels are $N \times N$ matrices that slide over the entire input image to extract the important features of an input image.
4. Pooling is required to reduce the dimensionality of the feature matrix. Sometimes, the feature matrix obtained after the convolutional layer results in a large number of features that need to be reduced so that it becomes computationally less expensive.
5. Padding is another type of layer present within CNN architecture. Here, additional zeroes are included to the edges of the matrix in order to read the edges precisely or get the output image similar to the one that was passed as the input image.
6. Once the matrix has been flattened by the pooling layer, the flattened vector is passed on to the fully connected layer. This layer is a neural network where weights are adjusted.

We now move on to the implementation of a TensorFlow neural network.

Importing CIFAR

The CIFAR-10 dataset is readily accessible in Python through the Keras library, which is part of TensorFlow, making it a convenient choice for developers and researchers working on machine learning projects, especially in image classification. In this article, we will explore CIFAR10 (classification of 10 image labels) from Keras/tensorflow.

What is the CIFAR10 Datasets?

The CIFAR-10 dataset contains 60,000 32×32 color images in 10 different classes, such as airplanes, cars, birds, cats, deer, dogs, frogs, horses, ships, and trucks.

CIFAR10 dataset consists of black and white images categorized into 10 types of clothing items, each represented by an integer label ranging from 0 to 9. This structure ensures clarity and organization in the data, facilitating effective classification tasks. CIFAR-10 dataset is a collection of images that are commonly used to train machine learning and computer vision algorithms. It is one of the most widely used datasets for machine learning research.

Full Form of CIFAR10 DataSet

The CIFAR-10 dataset stands for Canadian Institute For Advanced Research Dataset, where 10 stands for the count of representation classes, as discussed above.

Characteristics of CIFAR10 Dataset

The common characteristics of CIFAR10 dataset include:

- Number of Instances: 60,000 images
- Training Set:
 - ❑ 50,000 images
 - ❑ Each image is a 32×32 color image (RGB), resulting in a shape of (32, 32, 3).
 - ❑ Images are divided into 10 classes, with 5,000 images per class.
- Test Set:
 - ❑ 10,000 images
 - ❑ Same structure as the training set, with 1,000 images per class.
- Pixel Values: Each pixel value (0-255) represents the grayscale intensity of the corresponding pixel in the image.
- Target: Target Column represents the type of clothing item (0-9)
- Number of Attributes: 1 (32×32 pixels = 1024 pixels)

Structure of the CIFAR10 dataset:

(x_train, x_test): These variables contain the pixel data for the images.

x_train is the training set of the images, and

x_test is the testing set.

The images are 32×32 pixels in size and are represented as a numpy array of shape (32, 32, 3), where 3 stands for the three color channels (RGB).

(y_train, y_test): These are the corresponding labels for the images. Each label is an integer from 0 to 9, representing the class of representation, i.e.:

(Label) -> (Class)

0 -> Airplane

1 -> Automobile

2 -> Bird

3 -> Cat

4 -> Deer

5 -> Dog

6 -> Frog

7 -> Horse

8 -> Ship

9 -> Truck



R.M.K.
GROUP OF
INSTITUTIONS

How to Load CIFAR10 Datasets in Keras?

To load the CIFAR-10 dataset using Keras, you can use the CIFAR10 module from tensorflow.keras.datasets.

Syntax:

```
from tensorflow.keras.datasets import cifar10
```

```
# Load the CIFAR-10 dataset
```

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```


Significance of CIFAR10 in Machine Learning

The CIFAR-10 dataset holds significant importance in the field of machine learning for several reasons:

Benchmark Dataset: CIFAR-10 serves as a benchmark dataset for testing the performance of various machine learning algorithms, particularly in the domain of computer vision. Its popularity stems from its moderate size, making it suitable for experimentation and benchmarking without requiring extensive computational resources.

Real-World Image Classification: The CIFAR-10 dataset consists of 60,000 32×32 color images across 10 classes, with each class representing a different object category (e.g., airplane, automobile, bird, cat, etc.). This diversity makes CIFAR-10 a suitable dataset for training and evaluating image classification models on real-world, diverse image data.

Transfer Learning and Pre-Trained Models: CIFAR-10 is often used for transfer learning experiments, where models pre-trained on larger datasets (e.g., ImageNet) are fine-tuned on CIFAR-10 to adapt them to specific classification tasks. This approach leverages the learned representations from large-scale datasets to improve performance on smaller datasets like CIFAR-10.

Complexity: Despite its small size and relatively low resolution, CIFAR-10 remains a challenging dataset for machine learning models due to the variety of object classes, background clutter, and variations in object appearance and orientation within each class.

Applications of the CIFAR10 Dataset:

The CIFAR-10 dataset, with its collection of 60,000 images across 10 different classes (airplanes, cars, birds, cats, deer, dogs, frogs, horses, ships, and trucks), serves as a fundamental resource for various applications and research in the field of computer vision and machine learning. Here are some key applications and uses of the CIFAR-10 dataset:

Benchmarking Models: CIFAR-10 is widely used to benchmark the performance of image recognition algorithms and neural network architectures. It helps researchers and developers compare the efficacy of different models under consistent conditions.

Training Convolutional Neural Networks (CNNs): Due to its moderate complexity and size, CIFAR-10 is excellent for training CNNs from scratch. It allows for rapid experimentation with network architectures, hyperparameters, and training procedures without the computational expense required for larger datasets like ImageNet.

Pre-training for Transfer Learning: CIFAR-10 can be used for pre-training models that are then fine-tuned on more specialized or smaller datasets. This is particularly useful when computational resources are limited or when the target dataset is too small to train a deep network effectively from scratch.

Educational Purposes: CIFAR-10 is commonly used in academic courses and tutorials related to machine learning and computer vision. It is complex enough to teach nuanced concepts of deep learning, yet simple enough for educational use.

Feature Learning: Researchers use CIFAR-10 to develop and test algorithms for learning feature representations from images. These learned features can be crucial for tasks such as image retrieval, classification, and anomaly detection.

Development of New Algorithms: Beyond traditional image classification, CIFAR-10 is used to develop new types of learning algorithms, such as semi-supervised learning, unsupervised learning, and self-supervised learning methods.

Real-time Object Recognition: Models trained on CIFAR-10 can be adapted to work in real-time applications, such as video surveillance and autonomous vehicles, where recognizing objects quickly and accurately is critical.

The CIFAR-10 dataset, readily accessible through the Keras library in Python, is a cornerstone in the realm of machine learning and computer vision. With its collection of 60,000 32×32 color images across 10 distinct classes, CIFAR-10 serves as a fundamental resource for various applications and research endeavors.

ASSIGNMENT – UNIT II

1. Explain the purpose of creating image encodings in the context of deep learning.(K6)
2. Describe the two main components of the loss function in a Variational Autoencoder (VAE)(K5)
3. What is the role of Inverse Autoregressive Flow (IAF) in enhancing the expressiveness of the variational posterior in VAEs.(K4)
4. How do you import and preprocess the CIFAR-10 dataset using TensorFlow 2(K6)
5. Outline the steps to build a simple Convolutional Neural Network (CNN) in TensorFlow 2 for encoding images.(K3)

PART A- UNIT-II

1.What is an image encoding in the context of deep learning?

Answer: An image encoding is a compact, lower-dimensional representation of an image that retains important features necessary for various tasks like classification, generation, or clustering.

2.Compare the effectiveness of image encodings with raw pixel data for image classification tasks.

Answer: Image encodings are generally more effective than raw pixel data because they abstract and condense the most important features, reducing noise and irrelevant information, which helps in improving the model's performance and training efficiency.

3.What are the two main components of the variational objective in VAEs?

Answer: The two main components are the reconstruction loss and the KL divergence loss.

4.Design an experiment to test the effect of different KL divergence weights on the performance of a VAE.

Answer: Train multiple VAEs on the same dataset with varying KL divergence weights. Evaluate and compare their reconstruction errors and the quality of generated samples using metrics like Mean Squared Error for reconstructions and visual inspection or FID scores for generated samples.

5.What is the purpose of Inverse Autoregressive Flow in variational inference?

Answer: The purpose of IAF is to enhance the expressiveness of the variational posterior by applying a sequence of invertible transformations to the latent variables, making the posterior distribution more flexible and capable of capturing complex dependencies.

6. Evaluate the benefits and potential drawbacks of using IAF in a VAE.

Answer: Benefits of using IAF include a more flexible and expressive posterior distribution, leading to better model performance in terms of reconstruction and generation. Potential drawbacks include increased computational complexity and training time, as well as the need for careful tuning of the model's hyperparameters.

7.What is the CIFAR-10 dataset?

Answer: The CIFAR-10 dataset is a collection of 60,000 32x32 color images in 10 different classes, with 6,000 images per class, commonly used for training machine learning and computer vision models.

8.What are the main components of a Convolutional Neural Network (CNN)?

Answer: The main components of a CNN are convolutional layers, pooling layers, and fully connected (dense) layers. Convolutional layers detect features, pooling layers reduce dimensionality, and dense layers perform classification or regression tasks..

9.Explain how pooling layers contribute to the functionality of a CNN.


Answer: Pooling layers contribute to the functionality of a CNN by reducing the spatial dimensions of the feature maps, which helps in lowering the computational load, reducing overfitting, and providing a form of spatial invariance to the model.

10.Analyze the impact of using different activation functions in the convolutional layers of a CNN.

Answer: Different activation functions impact the CNN's performance by affecting the learning capacity and non-linearity of the model. ReLU is commonly used due to its simplicity and effectiveness, but other functions like Leaky ReLU or ELU might provide better performance by addressing issues like dying neurons or vanishing gradients.

PART B- UNIT II

- 1.Explain why it is necessary to normalize pixel values when importing the CIFAR-10 dataset for deep learning tasks.(K3)
- 2.Compare the computational advantages and disadvantages of using different activation functions (e.g., ReLU, sigmoid, tanh) in the layers of a CNN designed for CIFAR-10 classification.(K4)
- 3.Describe how the loss function is constructed to balance reconstruction accuracy and latent space regularization.(K2)
- 4.Design a Variational Autoencoder (VAE) architecture with TensorFlow 2 that incorporates a variational objective.(K3)
- 5.Critically evaluate the role of Inverse Autoregressive Flow (IAF) in enhancing the flexibility of the approximate posterior distribution in a VAE. (K5)



SUPPORTIVE ONLINE COURSES – UNIT II

❁ <https://www.classcentral.com/course/image-compression-generation-vae-19624>

❁ <https://arxiv.org/abs/1606.04934>

❁ <https://medium.com/fenwicks/tutorial-2-94-accuracy-on-cifar10-in-2-minutes-7b5aaecd9cdd>

❁ <https://www.coursera.org/specializations/tensorflow2-deeplearning>



Real Time Applications in Day to Day life and to Industry

1. **Image Classification:** Image classification is a fundamental task in computer vision that involves training a model to recognize and categorize images into predefined classes or labels. TensorFlow provides a powerful platform for developing image classification models.

Salient Key Features:

- 1.Data Preparation
- 2.Pre-trained Models
- 3.Transfer Learning
- 4.Custom Model Development
- 5.Training and Optimization
- 6.Evaluation Metrics
- 7.Deployment
- 8.Applications



2. Speech Recognition: Speech recognition is a technology that converts spoken language into written text. TensorFlow provides powerful tools and models for developing speech recognition systems. Here are the key features:

1. Acoustic Modeling
2. Language Modeling
3. Automatic Speech Recognition (ASR)
4. Training Data Augmentation
5. Transfer Learning
6. Speech Enhancement
7. On-device Speech Recognition



3. Natural Language Processing (NLP) pertains to the domain of artificial intelligence that centers on the interaction between computers and human language. TensorFlow provides a diverse array of tools and models specifically designed to facilitate the development of applications related to NLP. The following are the principal characteristics:

1. Text Preprocessing
2. Word Embeddings
3. Sequence Modeling
4. Text Classification
5. Language Translation
6. Text Generation
7. Transfer Learning



PRESCRIBED TEXT BOOKS AND REFERENCE BOOKS

TEXT BOOKS:

1. Ben Auffarth, Generative AI with LangChain, Packt Publishing, 2023.
2. Amit Bahree, Generative AI in Action, Manning Publication, First Edition, 2023.

REFERENCES:

1. David Foster, Generative Deep Learning, 2nd Edition, O'Reilly Media, 2023.
2. Numa Dhamani and Maggie Engler, Introduction to Generative AI, Manning Publication, First Edition, 2024.
3. Valentina Alto, Modern Generative AI with ChatGPT and OpenAI Models, Packt publications, 2024.



R.M.K.
GROUP OF
INSTITUTIONS

MINI PROJECT SUGGESTIONS

❁**TEAM 1:** Detecting Spam using TensorFlow

❁**TEAM 2:** Image Classification with TensorFlow

❁**TEAM3:** Optical Character Recognition using TensorFlow

❁**TEAM 4:** Sudoku Solver using TensorFlow

❁**TEAM 5:** Optical Character Recognition using TensorFlow



TOUGH QUESTIONS SIMILAR TO GATE

1. In a VAE, the Evidence Lower Bound (ELBO) is given by:

$$\mathcal{L}(\theta, \phi) = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - \text{KL}(q_{\phi}(z|x) \| p(z))$$

What happens if the KL divergence term dominates the reconstruction loss?

- A) The VAE generates diverse but noisy samples.
- B) The VAE ignores latent structure, behaving like a plain autoencoder.
- C) The VAE collapses to the prior, ignoring input data.
- D) The VAE overfits the training data.

Ans: C

2. For two Gaussians

$p = \mathcal{N}(\mu_1, \sigma_1^2)$ and $q = \mathcal{N}(\mu_2, \sigma_2^2)$, the KL divergence $\text{KL}(p \| q)$ is:

- A) $\log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$
- B) $\log \frac{\sigma_1}{\sigma_2} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$
- C) $\log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_2^2 + (\mu_1 - \mu_2)^2}{2\sigma_1^2} - \frac{1}{2}$
- D) $\log \frac{\sigma_1}{\sigma_2} + \frac{\sigma_2^2 + (\mu_1 - \mu_2)^2}{2\sigma_1^2} - \frac{1}{2}$

Ans: C

TOUGH QUESTIONS SIMILAR TO GATE

3. How does IAF differ from MAF (Masked Autoregressive

Flow)?

$$\mathcal{L}(\theta, \phi) = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - \text{KL}(q_{\phi}(z|x) \| p(z))$$

- A) IAF is faster for sampling but slower for density estimation.
- B) IAF is faster for density estimation but slower for sampling.
- C) IAF cannot model multimodal distributions.
- D) IAF requires invertible neural networks.

Ans: A

4. For the reparameterization trick is $z \sim \mathcal{N}(\mu, \sigma^2)$,

- A) $z = \mu + \sigma \cdot \epsilon$ where $\epsilon \sim \mathcal{N}(0, 1)$.
- B) $z = \mu + \epsilon$ where $\epsilon \sim \mathcal{N}(0, \sigma^2)$.
- C) $z = \sigma + \mu \cdot \epsilon$ where $\epsilon \sim \mathcal{N}(0, 1)$.
- D) $z = \mu \cdot \sigma \cdot \epsilon$ where $\epsilon \sim \mathcal{N}(0, 1)$.

Ans: A

TOUGH QUESTIONS SIMILAR TO GATE

5. Which code correctly loads CIFAR-10 and normalizes pixels to

$[-1,1][-1,1]$?

A) `(x_train, y_train), _ = tf.keras.datasets.cifar10.load_data()`

`x_train = x_train / 255.0`

B) `(x_train, y_train), _ = tf.keras.datasets.cifar10.load_data()`

`x_train = (x_train / 127.5) - 1.0`

C) `(x_train, y_train), _ = tf.keras.datasets.cifar10.load_data()`

`x_train = (x_train / 127.5) - 1.0`

D) `(x_train, y_train), _ = tf.keras.datasets.cifar10.load_data()`

`x_train = tf.image.per_image_standardization(x_train) .`

Ans: B

6. Which code defines a VAE encoder with stochastic sampling?

A) `class Encoder(tf.keras.layers.Layer):`

`def call(self, x):`

`z_mean = dense(x)`

`z_log_var = dense(x)`

`return z_mean, z_log_var`

TOUGH QUESTIONS SIMILAR TO GATE

B) `class Encoder(tf.keras.layers.Layer):`

```
def call(self, x):
```

```
    z_mean = dense(x)
```

```
    z_log_var = dense(x)
```

```
    epsilon = tf.random.normal(tf.shape(z_mean))
```

```
    return z_mean + tf.exp(z_log_var) * epsilon
```

C) `class Encoder(tf.keras.layers.Layer):`

```
def call(self, x):
```

```
    z_mean = dense(x)
```

```
    z_log_var = dense(x)
```

```
    epsilon = tf.random.normal(tf.shape(z_mean))
```

```
    return z_mean + tf.exp(0.5 * z_log_var) * epsilon
```

D) `class Encoder(tf.keras.layers.Layer):`

```
def call(self, x):
```

```
    z = dense(x)
```

```
    return tf.random.normal(z)
```

Ans : C

TOUGH QUESTIONS SIMILAR TO GATE

7. The reconstruction loss for a VAE with binary cross-entropy is suitable for:

- A) Grayscale images
- B) RGB images normalized to $[0, 1]$
- C) RGB images normalized to $[-1, 1]$
- D) Text data

Ans : B

8. For a flow $z = f(x)$, the probability $p(x)$ is computed as:

- A) $p(x) = p(z) \cdot \det \left| \frac{\partial f}{\partial x} \right|$
- B) $p(x) = p(z) \cdot \left| \frac{\partial f}{\partial x} \right|^{-1}$
- C) $p(x) = p(z) \cdot \det \left| \frac{\partial f}{\partial x} \right|^{-1}$
- D) $p(x) = p(z) + \log \left| \frac{\partial f}{\partial x} \right|$

Ans : B

TOUGH QUESTIONS SIMILAR TO GATE

9. Your VAE generates blurry images. What is the least likely fix?

- A) Increase the latent space dimension.
- B) Replace MSE loss with SSIM loss.
- C) Use a stronger decoder.
- D) Remove the KL term.

Ans : D

10. How does IAF improve upon RealNVP?

- A) IAF scales to higher dimensions better.
- B) IAF avoids masking entirely.
- C) IAF uses autoregressive transforms in reverse.
- D) IAF eliminates the need for Jacobian calculations.

Ans : C

THANK YOU

Disclaimer:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

R.M.K

GROUP OF ENGINEERING INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS

R.M.K GROUP OF INSTITUTIONS



R.M.K
GROUP OF
INSTITUTIONS



Please read this disclaimer before proceeding:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.

Digital Course Material

22AI005

Introduction to Generative AI

Department : Information Technology

Batch/Year : 2022-2026/IV

Created by : Ms.G.K.Monica

Date : 20-08-2025

TABLE OF CONTENTS

S.NO	CONTENTS	SLIDE NO.
1	COURSE OBJECTIVES	6
2	PRE REQUISITES (COURSE NAMES WITH CODE)	7
3	SYLLABUS (WITH SUBJECT CODE, NAME, LTPC DETAILS)	8
4	COURSE OUTCOMES (6)	9
5	CO- PO/PSO MAPPING	10
6	LECTURE PLAN –UNIT 3	12
7	ACTIVITY BASED LEARNING –UNIT 3	14
8	TEST YOURSELF	15
9	LECTURE NOTES	16
10	ASSIGNMENT	50
11	PART A Q s (WITH K LEVEL AND CO)	51
12	PART B Q s (WITH K LEVEL AND CO)	53
13	SUPPORTIVE ONLINE CERTIFICATION COURSES	54
14	REAL TIME APPLICATIONS IN DAY TO DAY LIFE AND TO INDUSTRY	55
15	PRESCRIBED TEXT BOOKS & REFERENCE BOOKS	59
16	MINI PROJECT SUGGESTIONS	60

COURSE OBJECTIVES

The Course will enable learners to:

- ❖ To understand the basic concepts of Generative AI.
- ❖ To build Generative AI systems to generate images.
- ❖ To understand the concept used in Generative AI Models.
- ❖ To use various Generative AI models.
- ❖ To compare and use the various Large Language Models.
- ❖ To understand the basics of Prompt Engineering.



PRE REQUISITES

❁ PRE-REQUISITE CHART

**22AI005 INTRODUCTION TO
GENERATIVE AI**



OBJECTIVES:

- ❖ To understand the basic concepts of Generative AI.
- ❖ To build Generative AI systems to generate images.
- ❖ To understand the concept used in Generative AI Models.
- ❖ To use various Generative AI models.
- ❖ To compare and use the various Large Language Models.
- ❖ To understand the basics of Prompt Engineering.

UNIT I INTRODUCTION 9

Generative Models – Image transformation – Challenges - Deep Neural Networks – Perceptron – back propagation – CNN – RNN – Optimizer.

UNIT II IMAGE GENERATION 9

Creating encodings of images – variational objective – Inverse Autoregressive flow – Importing CIFAR – Creating the network from TensorFlow 2.

UNIT III GENERATIVE ADVERSARIAL NETWORKS 9

Generative Adversarial Networks – Vanilla GAN – Improved GANs – Progressive GAN – Challenges – Paired style transfer – Unpaired style transfer – Deepfakes – Modes of operation – key feature set – High level flow – Replacement – Re-enactment.

UNIT IV LARGE LANGUAGE MODELS 9

Overview of LLMs - Transformers – GPT – Types of LLMs – Key concepts – other Transformers – T5 – Generative Pre-Training Models – Multi-modal Models – DALL.E 2

UNIT V PROMPT ENGINEERING 9

Basics – In-Context Learning – In-Context Prompting – Techniques – Image Prompting – Prompt Hijacking – Challenges.

TOTAL: 45 PERIODS

COURSE OUTCOME

Course Code	Course Outcome Statement	Cognitive/ Affective Level of the Course Outcome	Expected Level of Attainment
Course Outcome Statements in Cognitive Domain			
C305.1	Elaborate the basic concepts of Generative AI	Understand K2	60%
C305.2	Build Generative AI systems to generate images	Analyse K4	60%
C305.3	Apply the concepts used in Generative AI Models	Apply K3	60%
C305.4	Use various Generative AI models.	Apply K3	60%
C305.5	Compare and use the various Large Language Models	Analyse K4	60%
C305.6	Analyze the basics of Prompt Engineering.	Apply K3	60%
Course Outcome Statements in Affective domain			
C305.7	Attend the classes regularly	Respond (A2)	95%
C305.8	Submit the Assignments regularly.	Respond (A2)	95%
C305.9	Participation in Seminar/Quiz/ Group Discussion/ Collaborative learning and content beyond syllabus	Valuing (A3)	95%

CO-PO/PSO MAPPING

Correlation Matrix of the Course Outcomes to Programme Outcomes and Programme Specific Outcomes Including Course Enrichment Activities

Course Outcomes (Cos)		Programme Outcomes (POs), Programme Specific Outcomes (PSOs)														
		PO 1	PO 2	PO 3	PO 4	PO 5	PO 6	PO 7	PO 8	PO 9	PO 10	PO 11	PO 12	PS O 1	PS O 2	PS O 3
		K 3	K 4	K 5	K 5	K 3 / K 5	A 2	A 3	A 3	A 3	A 3	A 3	A 2	K 3	K 3	K 3
C305.1	K 3	3	2	1	1	3						3		3	3	3
C305.2	K 4	3	3	2	2	3						3		3	3	3
C305.3	K 2	2	1									2		3	3	3
C305.4	K 3	3	2	1	1	3						3		3	3	3
C305.5	K 4	3	3	2	2	3						3		3	3	3
C305.6	K 3	3	2	1	1	3						3		2	2	2
C305.7	A 2												3			
C305.8	A 2								2	2	2		3			
C305.9	A 3						3	3		3	3		3			
C305		3	3	2	2	3	1	1	1	3	3	3	3	3	3	3

UNIT III GENERATIVE ADVERSARIAL NETWORKS

LECTURE PLAN – UNIT III

UNIT III GENERATIVE ADVERSARIAL NETWORKS							
Sl. No	TOPIC	NO OF PERIODS	PROPOSED LECTURE	ACTUAL LECTURE	PERTAINING CO(s)	TAXONOMY LEVEL	MODE OF DELIVERY
			PERIOD	PERIOD			
1	Generative Adversarial Networks	1			CO1	K2	PPT
2	Vanilla GAN – Improved GANs	1			CO1	K2	PPT
3	Progressive GAN – Challenges	1			CO1	K2	PPT
4	Paired style transfer – Unpaired style transfer	1			CO1	K3	PPT
5	Deepfakes	1			CO1	K3	PPT
6	Modes of operation – key feature set	1			CO1	K2	PPT
7	High level flow	1			CO1	K2	PPT
8	Replacement	1			CO1	K2	PPT
9	Re-enactment.	1			CO1	K2	PPT

LECTURE PLAN – UNIT III

ASSESSMENT COMPONENTS

- ❖ AC 1. Unit Test AC 2. Assignment
- ❖ AC 3. Course Seminar AC 4. Course Quiz
- ❖ AC 5. Case Study
- ❖ AC 6. Record Work
- ❖ AC 7. Lab / Mini Project
- ❖ AC 8. Lab Model Exam AC 9. Project Review

MODE OF DELEIVERY

- MD 1. Oral presentation MD
- 2. Tutorial
- MD 3. Seminar
- MD 4 Hands On MD 5. Videos MD 6. Field Visit



R.M.K.
GROUP OF
INSTITUTIONS

TEST YOURSELF

1. What is the primary objective of a Generative Adversarial Network (GAN)?

- a) Image classification
- b) Image generation
- c) Text summarization
- d) Text translation

2. What are the two main components of a GAN?

- a) Generator and encoder
- b) Discriminator and encoder
- c) Generator and discriminator
- d) Encoder and discriminator

3. What is the training process in a GAN called?

- a) Supervised learning
- b) Reinforcement learning
- c) Unsupervised learning
- d) Adversarial learning

4. Which loss function is commonly used in GANs?

- a) Cross-entropy loss
- b) Mean squared error loss
- c) Binary logistic loss
- d) Kullback-Leibler divergence

5. Which type of GAN is designed to generate samples conditioned on specific input information?

- a) Unconditional GAN
- b) Wasserstein GAN
- c) Progressive GAN
- d) Conditional GAN

3.1.GAN(Generative Adversarial Network)

GAN(Generative Adversarial Network) represents a cutting-edge approach to generative modeling within deep learning, often leveraging architectures like convolutional neural networks. The goal of generative modeling is to autonomously identify patterns in input data, enabling the model to produce new examples that feasibly resemble the original dataset.

What is a Generative Adversarial Network?

Generative Adversarial Networks (GANs) are a powerful class of neural networks that are used for an unsupervised learning. GANs are made up of two neural networks, **a discriminator and a generator**. They use adversarial training to produce artificial data that is identical to actual data.

- ❖ The Generator attempts to fool the Discriminator, which is tasked with accurately distinguishing between produced and genuine data, by producing random noise samples.
- ❖ Realistic, high-quality samples are produced as a result of this competitive interaction, which drives both networks toward advancement.
- ❖ GANs are proving to be highly versatile artificial intelligence tools, as evidenced by their extensive use in image synthesis, style transfer, and text-to-image synthesis.
- ❖ They have also revolutionized generative modeling.

Through adversarial training, these models engage in a competitive interplay until the generator becomes adept at creating realistic samples, fooling the discriminator approximately half the time. Generative Adversarial Networks (GANs) can be broken down into three parts:

- ❖ **Generative:** To learn a generative model, which describes how data is generated in terms of a probabilistic model.

- ❖ **Adversarial:** The word adversarial refers to setting one thing up against another. This means that, in the context of GANs, the generative result is compared with the actual images in the data set. A mechanism known as a discriminator is used to apply a model that attempts to distinguish between real and fake images.
- ❖ **Networks:** Use deep neural networks as artificial intelligence (AI) algorithms for training purposes.

3.1.1.Types of GANs

There are several types of GANs (Generative Adversarial Networks), each designed to address specific challenges or improve upon the basic GAN architecture. Here are the most common types

1. Vanilla GAN:

- ❖ This is the simplest type of GAN. Here, the Generator and the Discriminator are simple a basic [multi-layer perceptrons](#).
- ❖ In vanilla GAN, the algorithm is really simple, it tries to optimize the mathematical equation using [stochastic gradient descent](#).

2. Conditional GAN (CGAN):

[CGAN](#) can be described as a [deep learning](#) method in which **some conditional parameters are put into place**.

- ❖ In CGAN, an additional parameter 'y' is added to the Generator for generating the corresponding data.
- ❖ Labels are also put into the input to the Discriminator in order for the Discriminator to help distinguish the real data from the fake generated data.

3. Deep Convolutional GAN (DCGAN):

[DCGAN](#) is one of the most popular and also the most successful implementations of GAN. It is composed of [ConvNets](#) in place of [multi-layer perceptrons](#).

- ❖ The ConvNets are implemented without max pooling, which is in fact replaced by convolutional stride.
- ❖ Also, the layers are not fully connected.

4. **Laplacian Pyramid GAN (LAPGAN):**

- ❖ The LAPGAN is a linear invertible image representation consisting of a set of band-pass images, spaced an octave apart, plus a low-frequency residual.
- ❖ This approach uses multiple numbers of Generator and Discriminator networks and different levels of the Laplacian Pyramid.
- ❖ This approach is mainly used because it produces very high-quality images. The image is down-sampled at first at each layer of the pyramid and then it is again up-scaled at each layer in a backward pass where the image acquires some noise from the Conditional GAN at these layers until it reaches its original size.

5. **Super Resolution GAN (SRGAN):**

- ❖ SRGAN as the name suggests is a way of designing a GAN in which a deep neural network is used along with an adversarial network in order to produce higher-resolution images.
- ❖ This type of GAN is particularly useful in optimally up-scaling native low-resolution images to enhance their details minimizing errors while doing so.

3.1.2. Architecture of GANs

A Generative Adversarial Network (GAN) is composed of two primary parts, which are the Generator and the Discriminator.

Generator Model

- ❖ A key element responsible for creating fresh, accurate data in a Generative Adversarial Network (GAN) is the generator model.

- ❖ The generator takes random noise as input and converts it into complex data samples, such text or images. It is commonly depicted as a deep neural network.
- ❖ The training data's underlying distribution is captured by layers of learnable parameters in its design through training.
- ❖ The generator adjusts its output to produce samples that closely mimic real data as it is being trained by using backpropagation to fine-tune its parameters.
- ❖ The generator's ability to generate high-quality, varied samples that can fool the discriminator is what makes it successful.

Generator Loss

- ❖ The objective of the generator in a GAN is to produce synthetic samples that are realistic enough to fool the discriminator.
- ❖ The generator achieves this by minimizing its loss function J_G . The loss is minimized when the log probability is maximized, i.e., when the discriminator is highly likely to classify the generated samples as real.
- ❖ The following equation is given below:

$$J_G = -\frac{1}{m} \sum_{i=1}^m \log D(G(z_i))$$

Where,

- J_G measure how well the generator is fooling the discriminator.
- $\log D(G(z_i))$ represents log probability of the discriminator being correct for generated samples.
- The generator aims to minimize this loss, encouraging the production of samples that the discriminator classifies as real $\log D(G(z_i))$, close to 1.

Discriminator Model

- ❖ An artificial neural network called a discriminator model is used in Generative Adversarial Networks (GANs) to differentiate between generated and actual input.
- ❖ By evaluating input samples and allocating probability of authenticity, the discriminator functions as a binary classifier.
- ❖ Over time, the discriminator learns to differentiate between genuine data from the dataset and artificial samples created by the generator. This allows it to progressively hone its parameters and increase its level of proficiency.

Discriminator Loss

- ❖ The discriminator reduces the negative log likelihood of correctly classifying both produced and real samples.
- ❖ This loss incentivizes the discriminator to accurately categorize generated samples as fake and real samples with the following equation:

$$J_D = -\frac{1}{m} \sum_{i=1}^m \log D(x_i) - \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(z_i)))$$

- ❖ J_D assesses the discriminator's ability to discern between produced and actual samples.
- ❖ The log likelihood that the discriminator will accurately categorize real data is represented by $\log D(x_i)$
- ❖ The log chance that the discriminator would correctly categorize generated samples as fake is represented by $\log(1 - D(G(z_i)))$
- ❖ The discriminator aims to reduce this loss by accurately identifying artificial and real samples.


MinMax Loss

In a Generative Adversarial Network (GAN), the minimax loss formula is provided by:

$$\min_G \max_D (G, D) = E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p_z(z)} \left[\log (1 - D(g(z))) \right]$$

Where,

- ❖ G is generator network and D is the discriminator network
- ❖ Actual data samples obtained from the true data distribution $p_{data}(x)$ are represented by x.
- ❖ Random noise sampled from a previous distribution $p_z(z)$ (usually a normal or uniform distribution) is represented by z.
- ❖ $D(x)$ represents the discriminator's likelihood of correctly identifying actual data as real.
- ❖ $D(G(z))$ is the likelihood that the discriminator will identify generated data coming from the generator as authentic.



How does a GAN work?

The steps involved in how a GAN works:

Initialization:

Two neural networks are created: a Generator (G) and a Discriminator (D).

- ❖ G is tasked with creating new data, like images or text, that closely resembles real data.
- ❖ D acts as a critic, trying to distinguish between real data (from a training dataset) and the data generated by G.

Generator's First Move:

- ❖ G takes a random noise vector as input. This noise vector contains random values and acts as the starting point for G's creation process.
- ❖ Using its internal layers and learned patterns, G transforms the noise vector into a new data sample, like a generated image.

Discriminator's Turn:

D receives two kinds of inputs:

- ❖ Real data samples from the training dataset.
- ❖ The data samples generated by G in the previous step. D's job is to analyze each input and determine whether it's real data or something G cooked up. It outputs a probability score between 0 and 1. A score of 1 indicates the data is likely real, and 0 suggests it's fake.

The Learning Process:

Now, the adversarial part comes in:

- ❖ If D correctly identifies real data as real (score close to 1) and generated data as fake (score close to 0), both G and D are rewarded to a small degree. This is because they're both doing their jobs well.

- ❖ However, the key is to continuously improve. If D consistently identifies everything correctly, it won't learn much. So, the goal is for G to eventually trick D.

Generator's Improvement:

- ❖ When D mistakenly labels G's creation as real (score close to 1), it's a sign that G is on the right track. In this case, G receives a significant positive update, while D receives a penalty for being fooled.
- ❖ This feedback helps G improve its generation process to create more realistic data.

Discriminator's Adaptation:

- ❖ Conversely, if D correctly identifies G's fake data (score close to 0), but G receives no reward, D is further strengthened in its discrimination abilities.
- ❖ This ongoing duel between G and D refines both networks over time.

As training progresses, G gets better at generating realistic data, making it harder for D to tell the difference. Ideally, G becomes so adept that D can't reliably distinguish real from fake data. At this point, G is considered well-trained and can be used to generate new, realistic data samples.

3.1.3 Application Of Generative Adversarial Networks (GANs)

GANs, or Generative Adversarial Networks, have many uses in many different fields. Here are some of the widely recognized uses of GANs:

Image Synthesis and Generation : GANs are often used for picture synthesis and generation tasks, They may create fresh, lifelike pictures that mimic training data by learning the distribution that explains the dataset

Image-to-Image Translation : GANs may be used for problems involving image-to-image translation, where the objective is to convert an input picture from one domain to another while maintaining its key features. GANs may be used, for instance, to change pictures from day to night, transform drawings into realistic images, or change the creative style of an image.

Text-to-Image Synthesis : GANs have been used to create visuals from descriptions in text. GANs may produce pictures that translate to a description given a text input, such as a phrase or a caption. This application might have an impact on how realistic visual material is produced using text-based instructions.

Data Augmentation : GANs can augment present data and increase the robustness and generalizability of machine-learning models by creating synthetic data samples.

Data Generation for Training : GANs can enhance the resolution and quality of low-resolution images. By training on pairs of low-resolution and high-resolution images, GANs can generate high-resolution images from low-resolution inputs, enabling improved image quality in various applications such as medical imaging, satellite imaging, and video enhancement.

Advantages of GAN:

- Synthetic data generation
- High-quality results
- Unsupervised learning
- Versatility

Disadvantages of GAN

- Training Instability
- Computational Cost
- Overfitting
- Bias and Fairness
- Interpretability and Accountability

3.2.Improved GANs

Improved GANs refer to modifications of the original GAN architecture to address common issues such as **mode collapse**, **vanishing gradients**, and **instability** in training. Various GAN variants have been proposed to overcome these challenges:

1. Wasserstein GAN (WGAN):

- ❖ WGAN uses the Wasserstein distance to measure how close the generated data distribution is to the real data distribution.
- ❖ This modification improves training stability and prevents vanishing gradients.

2. Least Squares GAN (LSGAN):

- ❖ LSGAN replaces the binary cross-entropy loss with a least-squares objective, smoothing out the loss function and leading to more stable convergence.

3.Energy-based GAN (EBGAN):

- ❖ In EBGAN, the discriminator is treated as an energy function rather than a classifier, encouraging more stable and robust training.

Why Improvements Matter:

- ❖ Improves training convergence.
- ❖ Mitigates mode collapse (when the generator produces limited variations of output).
- ❖ Helps with optimization and preventing the discriminator from overpowering the generator.

3.3 Progressive GAN

- ❖ Progressive GANs are specifically designed to handle high-resolution image generation.
- ❖ The key innovation here is the progressive growing of layers during training.
- ❖ The process starts with low-resolution images, and new layers are gradually added to increase the image resolution step by step.
- ❖ By starting at a lower resolution, the GAN can focus on learning simple patterns first and then gradually refine its understanding as it deals with more complex patterns in higher resolutions.

Key Points:

- ❖ Used for generating high-resolution images (e.g., 1024x1024 or higher).
- ❖ Commonly applied in facial image generation.
- ❖ Training starts with low-res images, and higher-res layers are added progressively, leading to more stable training and better results.

3.4.Challenges in GANs

Training GANs is known to be difficult due to several key challenges:

Mode Collapse:

This happens when the generator produces very limited types of outputs (e.g., always generating similar images), ignoring the full diversity of the data distribution.

Training Instability:

GAN training requires a delicate balance between the generator and discriminator. If one becomes too strong compared to the other, training collapses.

Vanishing Gradients:

As the discriminator improves, it may output very confident predictions, causing gradients to vanish for the generator, making further training difficult.

Evaluation Metrics:

Measuring the quality of generated data is challenging. Common metrics include the Inception Score (IS) and the Fréchet Inception Distance (FID), but they are not perfect.

Why These Challenges Matter:

- ❖ Unstable training can prevent the GAN from learning effectively.
- ❖ Mode collapse results in lower diversity in generated data, which reduces the effectiveness of GANs for creative and generative tasks.

3.5.Paired Style Transfer

Paired Style Transfer refers to a technique in computer vision where an image's content is transformed by applying the style of another image, but with paired training data. In paired style transfer, the system is provided with two datasets: one containing the original images (content) and the other containing their corresponding styled versions (style).

This creates a supervised learning scenario, where each input image has a corresponding "styled" image, making it easier to train the model to learn the mapping between the two domains.

How Paired Style Transfer Works:

The process involves training a model (often using **convolutional neural networks**) to convert an image from the source domain (content image) into the target domain (styled image). Since the training data consists of paired examples, the model learns to understand how a particular style (e.g., Van Gogh painting style, black-and-white photo) transforms specific content (e.g., landscape photo, portrait).

Key Techniques:

Convolutional Neural Networks (CNNs):

The backbone of style transfer, CNNs can extract features from the input images and apply transformations that change the content image to match the style image.

Loss Functions:

These are used to ensure that the generated images resemble the desired outputs. In paired style transfer, both **content loss** (to preserve the content) and **style loss** (to apply the correct style) are calculated.

Image-to-Image Translation Networks:

Models like **Pix2Pix** are commonly used for paired style transfer. These networks take paired input-output data and learn the transformation using supervised learning.

Applications:

Artistic Style Transfer: Transforming photos into paintings by learning from pairs of photographs and corresponding paintings.

Photo Enhancement: Enhancing the quality of images by learning to map low-quality photos to high-quality ones.

Medical Imaging: Translating one type of medical image to another (e.g., CT scans to MRI images) by learning from paired medical data.

Examples:

Pix2Pix Model: One of the most popular architectures for paired style transfer, Pix2Pix uses a generator to create the styled image and a discriminator to evaluate whether the generated image is real or fake.

Example Task: Converting a sketch to a photorealistic image by training the model on paired data where sketches correspond to their real images.

Supervised Learning: Since the model is trained on paired data, it learns explicit mappings. For instance, converting night-time images to day-time images can be achieved by training on a dataset where each night-time image has a paired day-time equivalent.

Advantage:

High Accuracy: The paired data provides explicit guidance to the model, ensuring that the transformation between content and style is learned effectively.

Better Control: Since the style for each image is predefined, the outcome is more predictable and accurate.

Limitations:

Data Dependency: Paired style transfer requires a large amount of paired data, which can be difficult to obtain, especially for specific styles or domains.

Supervised Training: The need for supervision limits the flexibility of the model compared to unpaired methods, which can learn mappings without explicit pairs.

3.6.Unpaired Style Transfer

Unpaired Style Transfer refers to a technique in computer vision and image processing where the goal is to transfer the style of one image to another without having explicitly paired training data. In contrast to paired style transfer, where each input image has a corresponding styled version for training, unpaired style transfer operates with two separate datasets: one containing content images and the other containing style images. There are no direct correspondences between the two datasets.

Key Concept

- ❖The main challenge in unpaired style transfer is learning a mapping between two domains (e.g., photographs and paintings) without having access to pairs of corresponding images.
- ❖This requires the system to understand the underlying distribution of both domains and find a way to convert images from one domain to the other while preserving key content features.
- ❖One of the most popular models for unpaired style transfer is CycleGAN, which uses cycle consistency as a core idea to achieve this goal.

How Unpaired Style Transfer Works:

The process typically involves two mappings:

Mapping A→B: This takes an image from domain A (e.g., a photo) and transforms it into domain B (e.g., a painting).

Mapping B→A: This performs the reverse transformation, converting an image from domain B back into domain A.

However, since there are no direct pairs, the model must ensure that the transformation preserves the key content characteristics of the image. To achieve this, unpaired style transfer employs **cycle consistency loss**.

Key Techniques:

Cycle Consistency: This ensures that the transformations between domains are meaningful and reversible. ³²

Adversarial Loss: Both generators are trained adversarially, where each generator tries to fool the discriminator into believing that the generated images are real.

Content Preservation: The content of the original image is preserved while applying the stylistic features of the target domain.

Applications:

Unpaired style transfer is used in various creative and practical fields. Some popular applications include:

Artistic Style Transfer: Converting real photos into paintings without paired examples (e.g., transforming photographs into Van Gogh or Monet-like paintings).

Image Enhancement: Changing the weather or lighting conditions in images (e.g., converting daytime photos to nighttime scenes).

Domain Adaptation: Translating between different visual domains, such as converting horse images to zebra images, or turning sketches into more realistic images.

Advantage:

Flexibility

Generative Power

Adaptability

Limitations:

Quality Control

Training Complexity.

Mode Collapse

3.7. Deepfakes

Deepfakes are synthetic media—typically videos or images—where someone's likeness, voice, or actions are replaced or altered using deep learning techniques. The term "deepfake" comes from combining "deep learning" and "fake," highlighting the use of artificial intelligence to create highly realistic fake content.

Deepfakes primarily involve the **manipulation of videos, images, or audio** to make it appear as though a person is saying or doing something they did not actually do. While deepfakes can be used for legitimate purposes, such as entertainment or art, they are often associated with deceptive or malicious uses, such as misinformation, fraud, or defamation.

How Deepfakes Work

Deepfakes are typically created using **Generative Adversarial Networks (GANs)** or similar neural network architectures.

These networks are trained on large datasets of images or videos of the target person, learning the nuances of their appearance, expressions, and voice.

The key steps to create a deepfake include:

Data Collection:

A large set of images, video frames, or audio clips of the person whose likeness will be faked is gathered. This training data is critical for teaching the AI model how the person looks and sounds from various angles and in different contexts.

Training a GAN or Autoencoder:

- ❖ The collected data is fed into a deep learning model, often a GAN, or an **autoencoder**. The GAN architecture consists of two networks: a **generator** and a **discriminator**.³²
- ❖ The generator creates fake images or videos, while the discriminator evaluates their authenticity by comparing them to real images or videos. Through this adversarial process, the generator improves its ability to create realistic fakes.

Autoencoders are another common approach, often used for face swapping or replacement tasks. The encoder learns to compress an image into a lower-dimensional representation, and the decoder reconstructs the image from this representation. Once trained on two individuals, the decoder can swap faces in new images or videos.

Face or Voice Synthesis:

- ❖ Once trained, the model can generate synthetic media by blending or overlaying the learned representations of the target person onto another person in real-time.
- ❖ In videos, this often involves face-swapping, where the face of one person is replaced with the face of the target person, keeping the head movements and expressions synchronized.

Post-Processing:

- ❖ After generating the fake video or image, additional editing may be applied to improve the quality and smooth transitions between frames, ensuring the deepfake looks natural.
- ❖ This includes ensuring that shadows, lighting, and facial expressions match the surroundings.

Applications:

Deepfakes can be used across various domains, both for legitimate and harmful purposes. Some common applications include:

Entertainment and Art:

Deepfakes can be used to generate realistic visual effects in movies, creating younger or older versions of actors or digitally recreating deceased performers. It can also be used to animate historical figures or fictional characters for artistic purposes.

Social Media and Memes:

Deepfake technology has been used to create humorous or satirical content, where individuals' faces are swapped or altered for comedic effect. This is often seen in viral videos, where celebrities' faces are swapped into famous movie scenes or other public content.

Personalized Content:

Deepfakes can enable personalized avatars or digital assistants that look and sound like the user or another person. It has also been explored for creating personalized media, such as custom-tailored video messages featuring a celebrity's likeness.

Education and Training:

Deepfakes can be used for educational purposes, such as recreating historical figures or scenarios for teaching. For example, a historical figure's likeness could be used in a digital lecture to give students a more engaging learning experience.

3.8. Modes of Operation (GAN)

- ❖ **Modes of Operation** in the context of **Generative Adversarial Networks (GANs)** refer to the different ways GANs can be structured, modified, or extended to achieve specific tasks or improvements over the traditional Vanilla GAN model.
- ❖ Over time, several variants of GANs have been proposed to tackle limitations, improve performance, or adapt the model for particular use cases. Below is a detailed explanation of some key **GAN Variants and their Modes of Operation**:

1. Vanilla GAN (Basic GAN)

The Vanilla GAN is the original version of GANs, introduced by Ian Goodfellow in 2014. It consists of two neural networks: a generator and a discriminator, which are trained in an adversarial setting.

Generator: Learns to create realistic data from random noise.

Discriminator: Learns to distinguish between real data and fake data generated by the generator.

Mode of Operation:

The two networks are trained simultaneously. The generator tries to produce data that is indistinguishable from real data, while the discriminator tries to correctly classify real and generated data. The generator's goal is to fool the discriminator, while the discriminator's goal is to correctly identify real versus fake data.

Limitations:

Mode collapse: The generator may produce a limited variety of outputs.

Training instability: The training process is sensitive and can be difficult to balance, leading to convergence issues.

2. DCGAN (Deep Convolutional GAN)

Deep Convolutional GANs (DCGANs) are an extension of Vanilla GANs that use convolutional layers instead of fully connected layers, making them especially effective for generating images.

Mode of Operation:

DCGANs replace the fully connected layers with convolutional and transposed convolutional layers. The discriminator becomes a convolutional neural network (CNN) that classifies whether an image is real or generated. The generator uses transposed convolutional layers to create images from latent vectors (random noise).

Applications:

DCGANs are widely used for tasks such as image synthesis, super-resolution, and in-painting.

3. Conditional GAN (cGAN)

Conditional GANs extend the basic GAN framework by conditioning both the generator and discriminator on additional information (e.g., class labels or attributes). This allows control over the generated data.

Mode of Operation:

In a cGAN, the generator and discriminator both receive a condition (such as a label or a specific class) as input along with the random noise vector (for the generator) or the real/generated data (for the discriminator). The condition could be a class label, an attribute, or even another image.

Applications:

Image-to-image translation: For example, translating sketches into colored images.

Text-to-image synthesis: Generating images based on descriptive text input.

4. CycleGAN

CycleGAN is designed for unpaired image-to-image translation, where the task is to translate an image from one domain to another without having paired examples for training (e.g., converting photographs to paintings).

Mode of Operation:

CycleGAN introduces cycle consistency loss. This means that if you translate an image from domain A to domain B and then back to domain A, the resulting image should be similar to the original image. The cycle consistency ensures that important content is preserved during translation.

Applications:

Style transfer: Converting images between different artistic styles.

Cross-domain translation: Converting between different domains like photos ↔ paintings, horses ↔ zebras.

5. Progressive Growing GAN (ProGAN)

Progressive Growing GANs (ProGANs) incrementally grow the generator and discriminator architectures during training, starting with low-resolution images and progressively increasing the resolution.

Mode of Operation:

Training starts with small, low-resolution images, allowing the model to learn a rough approximation. As training progresses, new layers are added to both the generator and discriminator, allowing the model to gradually handle higher-resolution images. This reduces training instability and improves the quality of high-resolution outputs.

Applications:

Image generation: Used to create photorealistic high-resolution images.

Art creation: Used in generative art, including the creation of high-quality synthetic faces.

3.9.Key Feature Set

Generative Adversarial Networks have revolutionized the field of machine learning and AI, especially in tasks like image generation, data augmentation, and deepfakes. Here are some of the **key features** of GANs:

Adversarial Training: GANs involve two networks—a generator and a discriminator—that are trained in opposition to one another. This adversarial setup is the core idea behind GANs and leads to the generator improving as it tries to fool the discriminator, and vice versa.

Generative Modeling: GANs learn to model the distribution of data and generate new samples that resemble the training data. This ability to generate new data makes GANs unique compared to traditional machine learning models that primarily focus on classification or regression tasks.

Non-supervised Learning: GANs typically do not require labeled data (although some variants like conditional GANs may use labels). This makes them highly valuable in scenarios where labeled data is scarce or expensive to obtain.

High-Quality Synthetic Data: GANs can produce high-quality synthetic data, whether it be images, videos, or even text. For example, GANs are used to generate realistic human faces, artwork, and even music.

Flexibility: GANs are highly flexible and have been adapted to many different tasks. Variants like CycleGAN, Pix2Pix, and StyleGAN offer specialized solutions for tasks like image translation, style transfer, and high-resolution image synthesis.

Application Versatility: GANs have broad applications in various fields, such as synthetic data generation, data augmentation for machine learning models, image super-resolution, and even medical imaging.

3.10.High-Level Flow of GANs

The **high-level flow** of a GAN can be broken down into the interaction between the **generator** and **discriminator** in an adversarial setting. Here's a step-by-step breakdown:

Noise Input: The process starts with the generator receiving random noise as input. This noise is usually sampled from a latent space (often a uniform or normal distribution). The generator's task is to transform this random noise into a data sample that resembles real data.

Data Generation: The generator processes the noise input through its network layers to produce synthetic data. For instance, in image generation tasks, this could be a synthetic image.

Discriminator's Role: The discriminator, which is a binary classifier, receives both the real data from the training dataset and the synthetic data from the generator. It aims to classify whether each data point is real or generated (fake).

Feedback to Generator: The generator and discriminator are connected in a loop through the loss function. If the discriminator successfully identifies the synthetic data as fake, it provides a gradient update to the generator, effectively teaching it how to produce more realistic samples. If the discriminator fails to correctly classify the data as fake, it updates its own parameters to improve its ability to distinguish between real and fake data.

Adversarial Learning: This adversarial process continues until the generator becomes skilled enough to produce data that the discriminator cannot reliably distinguish from real data. In theory, training reaches a point where the discriminator is no better than random guessing, and the generator produces highly realistic outputs.

Convergence: The GAN training process is considered successful when the generator produces samples that are indistinguishable from the real data, as determined by the discriminator.

2.11. Replacement (GAN in Face Replacement or Object Replacement)

- ❖ Replacement refers to the use of GANs to substitute one element in an image or video with another, particularly for tasks like face swapping, object replacement, or deepfakes.
- ❖ The idea is to replace a specific part of an image (or multiple frames in a video) while maintaining other elements constant, such as the background, lighting, or context.

How GANs are used in Replacement:

Face Replacement: GANs can be trained to swap faces in images or videos. The generator network learns to synthesize a realistic face replacement by blending the new face onto the body and adjusting details like lighting and pose. Popular GAN architectures for this purpose include **StyleGAN** and **DeepFaceLab GAN**.

- ❖ **Training:** GANs for face replacement are typically trained on large datasets of the two faces to be swapped. They learn to map the facial features of one individual onto the face of another while maintaining expressions and orientations.
- ❖ **Real-Time Processing:** Recent advancements allow for real-time face replacement in videos. By using GANs, the face swap can be applied frame-by-frame, ensuring consistent and realistic transitions throughout the video.

Object Replacement: Beyond faces, GANs can be used for replacing or modifying objects in images. For example, replacing one car with another in a street scene or changing the appearance of an item in a product photo.

Challenges: Ensuring realism in replacements can be challenging, especially with varying lighting conditions, occlusions, or intricate textures. GANs must handle these variations to produce convincing replacements.

2.12. Re-enactment (GANs for Face and Body Re-enactment)

Re-enactment involves using GANs to manipulate or recreate specific movements, expressions, or gestures in an image or video, typically involving human faces or bodies. The goal is to map the movements or expressions of a source individual onto a target individual in a highly realistic manner.

How GANs are used in Re-enactment:

Face Re-enactment: GAN-based face re-enactment takes facial expressions, mouth movements, or eye gestures from a source video and transfers them onto the face of a target individual. For example, by analyzing a person's facial movements in a video, GANs can generate a new video where a different person exhibits the same expressions and lip movements, mimicking speech or reactions.

- ❖ **Training Process:** GANs are trained on datasets consisting of videos or images of the target individual. During training, the generator learns to re-enact the target's facial expressions based on the source input. The discriminator, in turn, learns to distinguish between real re-enactments and synthetic ones, improving the generator's performance over time.
- ❖ **Application in Deepfakes:** In the context of deepfakes, re-enactment is often used to create fake videos of a person saying or doing things they did not actually say or do. For instance, re-enactment techniques can simulate a famous person speaking with the lip movements corresponding to any given audio track.

Body Re-enactment: Similar to face re-enactment, GANs can also be used for full-body re-enactment, where a person's body movements and gestures are replicated onto another person. This is especially useful in virtual reality, gaming, or motion capture for movies. ³²

Challenges: Realistic re-enactment requires detailed modeling of how the face or body behaves in different situations. This involves capturing minute changes in expressions, lighting, and camera angles. GANs must also handle potential occlusions, such as hair or glasses covering parts of the face.

Applications:

- 1. Entertainment:** Re-enactment GANs are used in films and media to recreate actors' performances or animate historical figures.
- 2. Virtual Reality and Gaming:** Motion capture and re-enactment GANs can make virtual avatars behave more realistically by mapping real human expressions and movements onto digital characters.
- 3. Human-Computer Interaction:** GANs can be used to develop more interactive and lifelike digital assistants by replicating human expressions and behaviors.

ASSIGNMENT – UNIT III

1. What improvements have been introduced in GAN architectures to tackle the issues of instability and mode collapse? (k3)
2. How does the layer-wise progressive growth of Progressive GANs contribute to generating high-quality, high-resolution images? What training challenges does this approach address effectively?(k4)
3. What are the common challenges in training GANs, including mode collapse and vanishing gradients? Discuss techniques used in recent research to overcome these difficulties.(k5)
4. What is the difference between Paired Style Transfer (like Pix2Pix) and Unpaired Style Transfer (like CycleGAN)? (k4)
5. How are GANs utilized to generate deepfakes, specifically in replacement and re-enactment modes? (k5)

PART A- UNIT-III

1. What is the basic architecture of a Vanilla GAN?

A Vanilla GAN consists of two neural networks: a generator, which creates fake data, and a discriminator, which classifies data as real or fake. They are trained together in a minimax game until the generator produces realistic data that the discriminator cannot distinguish from real data.

2. How does Wasserstein GAN (WGAN) differ from Vanilla GAN?

WGAN uses the Wasserstein distance (earth mover's distance) as a loss function instead of binary cross-entropy. This helps stabilize training and mitigates problems like mode collapse by providing smoother gradients for the generator.

3. What is mode collapse in GANs?

Mode collapse occurs when the generator produces limited variations of the output, ignoring the diversity in the real data distribution. This results in a lack of variety in the generated samples.

4. How does Least Squares GAN (LSGAN) improve over Vanilla GAN?

LSGAN replaces the binary cross-entropy loss with a least-squares loss, which penalizes the difference between real and generated data more smoothly, improving stability and helping to avoid vanishing gradients.

5. What is Progressive GAN, and how does it work?

Progressive GAN generates high-resolution images by gradually increasing the resolution of the generator and discriminator layers during training. This progressive approach stabilizes training and improves the quality of the generated images.

.

PART A- UNIT-III

6. What are the benefits of progressively growing GANs?

Progressive GANs reduce training instability and allow for the generation of high-resolution images, addressing issues of poor quality often seen in standard GANs when trained directly at high resolutions.

7. What is Paired Style Transfer?

Paired Style Transfer is used when there are corresponding pairs of images in two domains (e.g., an image and its sketch). Models like Pix2Pix are trained to learn direct mappings from one domain to another.

8. What is Unpaired Style Transfer?

Unpaired Style Transfer works when corresponding pairs of images from two domains are not available. CycleGAN is a model that uses cycle consistency loss to learn mappings between unpaired data, such as photos from two different styles or environments.

9. What is the main challenge addressed by CycleGAN in style transfer?

CycleGAN addresses the lack of paired training data by using cycle consistency loss, ensuring that when an image is transferred from one domain to another and back again, it remains consistent with the original image.

10. How do GANs contribute to the creation of deepfakes?

GANs are used to generate deepfakes by training the generator to create realistic facial or object images. These fake images can then be swapped (replacement) or manipulated (re-enactment) to produce highly realistic videos or photos.

11. What are the ethical concerns surrounding deepfakes?

Deepfakes can be used for misinformation, fraud, and defamation, posing significant ethical challenges. They raise concerns about privacy, consent, and the potential for malicious use in creating misleading media content.

PART A- UNIT-III

12. What is the difference between replacement and re-enactment in deepfakes?

Replacement involves substituting one face with another in an image or video, while re-enactment involves manipulating a person's facial expressions or movements in a video to create new, artificial actions.

13. What role does the discriminator play in GANs?

The discriminator evaluates the data it receives, classifying it as real or fake. Its feedback helps the generator improve by penalizing it when the generated data does not closely resemble the real data.

14. What is the generator's objective in a GAN?

The generator's objective is to create data that is so realistic that the discriminator cannot distinguish it from the actual data, effectively "fooling" the discriminator.

15. How does WGAN help mitigate vanishing gradients in GAN training?

WGAN uses a loss function based on Wasserstein distance, which provides a more stable training signal and prevents vanishing gradients, thus allowing the generator to learn effectively even when the discriminator becomes very strong.

16. What is the purpose of cycle consistency loss in CycleGAN?

Cycle consistency loss ensures that when an image is transferred from one domain to another and back again, the output is the same as the original input. This enforces consistency and prevents the generator from creating arbitrary mappings.

17. How does Progressive GAN handle the challenge of training at high resolutions?

Progressive GANs start training at a low resolution and gradually add higher-resolution layers to both the generator and discriminator. This progressive growth helps avoid the instability

PART A- UNIT-III

and poor performance often encountered when training directly at high resolutions.

18. What are some common applications of Paired Style Transfer?

Paired Style Transfer is commonly used in applications such as image-to-image translation, for example, converting sketches to photos, black-and-white images to color, or day-to-night transformations in photography.

19. How are GANs used in style transfer tasks?

GANs are used in style transfer to learn the mapping between two domains, such as transferring artistic styles to photographs or converting images from one domain (e.g., summer) to another (e.g., winter).

20. What are the main modes of operation in deepfakes?

The two primary modes of operation in deepfakes are replacement, where one person's face is swapped with another's, and re-enactment, where a person's expressions or movements are modified to produce new behaviors in a video.

PART B- UNIT III

1. How does the progressive growth of layers in Progressive GANs enhance training stability and image quality? (k2)
2. Explain the use of Wasserstein distance in WGAN. How does it improve training stability and address mode collapse compared to the original GAN loss function?(k3)
3. How do Least Squares GAN (LSGAN) and Wasserstein GAN (WGAN) differ in their approach to stabilizing GAN training? What are the effects of these improvements on the diversity and quality of generated samples?(k3)
4. What are the common challenges in GAN training, such as mode collapse and vanishing gradients? How do techniques like WGAN and LSGAN attempt to solve these issues?(k5)
5. How does CycleGAN perform Unpaired Style Transfer using cycle consistency loss? What challenges and limitations does this approach encounter in handling unpaired data?(k5)
6. How do Progressive GANs enable high-resolution image generation while maintaining stability? Discuss how the progressive addition of layers contributes to the final output's detail and quality.(k2)

SUPPORTIVE ONLINE COURSES – UNIT III


- <https://www.coursera.org/specializations/deep-learning>
- <https://www.coursera.org/specializations/generative-adversarial-networks-gans>
- <https://www.datacamp.com/>



Real Time Applications in Day to Day life and to Industry

1.Object detection:

CNN has been applied to object recognition across images by classifying objects based on shapes and patterns found within an image. CNN models have been created that can detect a wide range of objects from everyday items such as food, celebrities, or animals to more unusual ones including dollar bills and guns. Object detection is performed using techniques such as semantic or instance segmentation. CNNs have been used to localize and identify objects within images as well as create different views of those objects such as for use in drones or self-driving cars.




2.Speech Recognition

RNNs can be used for predicting phonetic segments considering sound waves from a medium as an input source .The set of inputs consists of phoneme or acoustic signals from an audio which are processed in a proper manner and taken as inputs. The RNN network will compute the phonemes and then produce a phonetic segment along with the likelihood of output.The steps used in speech recognition are as follows:-


The input data is first processed and recognized through a neural network. The result consists of a varied collection of input sound waves.

The information contained in the sound wave is further classified by intent and through keywords related to the query.



3. Language Modelling and Generating Text

Taking a sequence of words as input, we try to predict the possibility of the next word. This can be considered to be one of the most useful approaches for translation since the most likely sentence would be the one that is correct. In this method, the probability of the output of a particular time-step is used to sample the words in the next iteration.



PREScribed TEXT BOOKS AND REFERENCE BOOKS

TEXT BOOKS:

1. "GANs in Action: Deep Learning with Generative Adversarial Networks" by Jakub Langr and Vladimir Bok
2. "Deep Learning with Python" by François Chollet
3. "Deep Learning with TensorFlow and Keras" by Amita Kapoor

REFERENCES:

1. Goodfellow, Ian, et al. ["Generative Adversarial Nets"](#), 2014.
2. Karras, Tero, et al. ["Progressive Growing of GANs for Improved Quality, Stability, and Variation"](#), 2018.
3. Isola, Phillip, et al. ["Image-to-Image Translation with Conditional Adversarial Networks"](#), 2017.

MINI PROJECT SUGGESTIONS

- ✓ **TEAM 1:** Build a simple Vanilla GAN from scratch to generate images (e.g., handwritten digits using the MNIST dataset).
- ✓ **TEAM 2:** Progressive Growing of GANs for High-Resolution Image Generation
- ✓ **TEAM 3:** Paired and Unpaired Image-to-Image Translation using GANs (Style Transfer)
- ✓ **TEAM 4:** Deepfake Creation and Detection
- ✓ **TEAM 5:** GANs for Image Re-enactment and Replacement



THANK YOU

Disclaimer:

This document is confidential and intended solely for the educational purpose of RMK Group of Educational Institutions. If you have received this document through email in error, please notify the system manager. This document contains proprietary information and is intended only to the respective group / learning community as intended. If you are not the addressee you should not disseminate, distribute or copy through e-mail. Please notify the sender immediately by e-mail if you have received this document by mistake and delete this document from your system. If you are not the intended recipient you are notified that disclosing, copying, distributing or taking any action in reliance on the contents of this information is strictly prohibited.